SANKHYA AI RESEARCH · KATHMANDU · 2025

The Adoption Paradox

How Kathmandu's professionals use AI faster than they understand it — and faster than anyone is governing it

A cross-sector study of artificial intelligence awareness, adoption, trust and governance readiness across eight sectors in Kathmandu, Nepal.

334	8	98%	22%
PROFESSIONALS SURVEYED	INDUSTRY SECTORS	HAVE HEARD OF AI	UNDERSTAND IT AND USE IT DAILY

Sankhya AI — Your Intelligence Partner
Kathmandu, Nepal

Fieldwork March–April 2025
Published 2025

PUBLICATION DETAILS

The Adoption Paradox

How Kathmandu's professionals use AI faster than they understand it — and faster than anyone is governing it. A cross-sector study of artificial intelligence awareness, adoption, trust and governance readiness across eight sectors in Kathmandu, Nepal.

PUBLISHED BY

Sankhya AI
Your Intelligence Partner
Putalisadak, Kathmandu
Nepal

EDITION AND DATE

First edition, 2025.
Fieldwork conducted March–April 2025, Kathmandu district, Nepal.

AUTHORS

Indra Giri, Chief Executive Officer, Sankhya AI
Dr Ram Narayan Shrestha, Assistant Professor, Kathmandu University
Study design, analysis and authorship. Reviewed and endorsed by the Editorial Board listed on page 4; full acknowledgements on page 5.

RIGHTS

HOW TO CITE THIS REPORT

Giri, I. and Shrestha, R. N. (2025). *The Adoption Paradox: AI awareness, adoption and governance readiness across eight sectors in Kathmandu, Nepal.* Kathmandu: Sankhya AI.

```
@techreport{giri2025adoption,
  title = {The Adoption Paradox},
  author = {Giri, Indra and Shrestha, Ram
Narayan},
  year = {2025},
  institution = {Sankhya AI},
  address= {Kathmandu, Nepal},
  type = {Research report}
}
```

DISCLAIMER

The findings, interpretations and conclusions in this report are those of the authors and do not necessarily represent the views of the individuals named in the Editorial Board or of the institutions with which any of them are affiliated. Board members contributed in a personal capacity.

This report is based on a purposive sample of 334 respondents in Kathmandu district. It is not a probability sample and its findings should not be projected to population totals for Nepal or for Kathmandu. Confidence intervals are not reported. Limitations are set out in full in Appendix A.

Nothing in this report constitutes legal, financial, clinical or investment advice. Sector-specific statements describe what respondents reported, not what any organisation should do.

LANGUAGE

The survey instrument was administered in Nepali and English. All figures in this report are derived from the English-language export of the response file; response labels were restored to the wording of the English questionnaire, as documented in Appendix A.

MESSAGE FROM THE CHIEF EXECUTIVE OFFICER

We measured what everyone assumed they already knew

Nepal has spent two years being told that artificial intelligence is coming. This report is the first attempt to measure how much of it has already arrived — and to say honestly where it has not.

When we began designing this study in early 2025, the conversation about AI in Nepal was running almost entirely on anecdote. Everyone knew a student who used ChatGPT to write an essay. Everyone had seen a deepfake circulating on social media. Nobody could say, with evidence, how many professionals were actually using these tools, in which sectors, how often, or whether anyone understood what they were using.

That absence of evidence matters more than it might appear. Policy written without it will be written for an imagined country. Products built without it will be built for an imagined customer. Training programmes designed without it will teach the wrong people the wrong things. We commissioned this survey because we believed the gap between what Nepal assumed about AI and what was actually happening had become wide enough to be dangerous.

It turned out to be wider than we expected. Ninety-eight per cent of the professionals we interviewed have heard of artificial intelligence. Eighty-one per cent use it. But only twenty-two per cent both understand it well and use it every day. Between those two numbers sits the entire substance of this report.

Three findings deserve to be read carefully.

The first is that AI reached Nepal through the phone, not the workplace. Eighty-seven per cent of our respondents first learned about AI through social media. Only eight and a half per cent learned about it at work. No general-purpose technology of this consequence has ever entered a Nepali workplace this way before, and almost every other finding in this report is downstream of that single fact.

The second is that trust has run ahead of understanding. More people trust AI to help with important decisions than claim to understand what it does. In healthcare and legal practice — the two sectors where a wrong answer costs the most — we found people using these tools daily with no institutional guidance of any kind.

The third is the one that should concern all of us. Ninety-two per cent of the people we spoke to want artificial intelligence regulated. Fewer than half know that their own government has already drafted a policy to do it. That is not a failure of political will. It is a failure of communication, and it is fixable this year.

I want to be equally clear about what this report does not do. It does not measure whether AI has made anyone in Nepal more productive — we asked what people use, not what they gained. It covers Kathmandu, not the country. And it relies on people's assessment of their own competence, which is generous everywhere in the world. We have set these limitations out plainly in Appendix A rather than in a footnote, because a research organisation that hides its constraints cannot be trusted with its conclusions.

What we hope this report does is move the conversation from assertion to evidence. Nepal does not need to be persuaded that AI matters; ninety-eight per cent already know. What it needs now is to decide who is responsible for making it work, who is responsible for making it safe, and who is going to be left behind while that is decided.

Those are political and commercial questions, not technical ones. This report is our contribution to answering them well.

Indra Giri

Indra Giri

Chief Executive Officer, Sankhya AI
Kathmandu, 2025

EDITORIAL BOARD

Independent review and endorsement

The Editorial Board reviewed the study design, the analytical method and the findings presented in this report, and endorsed them for publication. Members served in a personal capacity and independently of the authors; their institutional affiliations are given for identification only.

Dr Krishna Sharma

Postdoctoral Research Fellow

Stanford University, United States

Reviewed the statistical treatment: the choice to report effect sizes and Monte Carlo p-values rather than asymptotic chi-square alone, the correction for multiple testing, the handling of variable response bases, and the decision not to post-stratify a purposive sample.

Mr Bishu Giri

Consultant

The World Bank, Washington DC, United States

Reviewed the policy chapter and the recommendations, and advised on framing the governance gap in terms that are actionable for Nepali public institutions and comparable with international AI-readiness assessments.

Dr Prabal Sharma

AI Engineer

University of Nebraska, United States

Reviewed the technical content of the tool-awareness and IT-sector sections, and advised on the distinction drawn throughout the report between consumption of AI systems and capability to build them.

SCOPE OF THE REVIEW

Board members reviewed the questionnaire, the cleaned dataset, the analytical code and the full draft report. Their endorsement covers the soundness of the method and the fidelity of the findings to the data. It does not extend to the interpretive material in Chapter 8 or the commercial judgements in Chapter 9, which are the authors' own, and it does not imply institutional endorsement by Stanford University, the World Bank or the University of Nebraska. Dr Ram Narayan Shrestha is an author of this report and therefore did not serve on the reviewing board.

WHAT WAS REVIEWED

The bilingual survey instrument; the cleaned and recoded dataset together with the record-level cleaning decisions; the analysis code, including index construction and every statistical test; and the complete draft report with all 28 exhibits.

WHAT CHANGED AS A RESULT

Two substantive revisions followed from the review. The Sankhya AI Readiness Index was rebuilt to weight use intensity rather than treat adoption as binary, and the treatment of the study's limitations was moved from a footnote into a full appendix with a dedicated statement on what the sample cannot support.

INDEPENDENCE

Board members received no payment for their review and hold no financial interest in Sankhya AI. They were given access to the underlying data and were free to withhold endorsement. Their comments were addressed in full before publication.

STUDY TEAM AND ACKNOWLEDGEMENTS

Who built this report

A survey of this kind is field labour before it is analysis. The quality of every figure in this report rests on the three enumerators who conducted the interviews.

AUTHORS

Indra Giri
Chief Executive Officer, Sankhya AI — study direction, analysis and authorship

Dr Ram Narayan Shrestha
Assistant Professor, Kathmandu University — study design, methodology and authorship

RESEARCH SUPPORT

Sudip Giri
Research Assistant — instrument translation, field coordination, data checking and coding support

DESIGN AND TECHNOLOGY

Priya Puri Giri
Report design and IT support — visual design, production and digital delivery

DATA COLLECTION

343 contacts were made face to face, in Nepali or English at the respondent's preference, between March and April 2025, of which **334 produced valid interviews**. Nine did not proceed past the consent stage.

Purnima Acharya
Field enumerator

Kamal Bahadur Rana Kshetri
Field enumerator

Narayan Wagle
Field enumerator

Enumerators were trained on the instrument, on informed consent and on neutral administration of attitudinal questions before fieldwork began.

Full review scope on page 4. Dr Ram Narayan Shrestha is an author and did not serve on the reviewing board.

ACKNOWLEDGEMENTS

Our first and largest debt is to the **334 professionals** who gave their time to this study. They were asked, without notice and in the middle of a working day, to answer thirty-one questions about a technology many of them were still forming a view on. Several spoke candidly about using AI in ways their employers did not know about. That candour is what makes this report worth reading, and we have protected it: no respondent is identifiable anywhere in this document or in any dataset we release.

We thank the schools, hospitals, banks, law firms, technology companies, newsrooms, hotels and workshops across Kathmandu that allowed our enumerators access to their staff, and the managers who agreed to the study without asking to see the questions first.

We are grateful to the members of the Editorial Board, who reviewed this work across three time zones and without payment, and whose criticism materially improved the index construction, the statistical treatment and the report's account of its own limitations.

ABOUT THIS REPORT

Contents and how to read this document

This report converts 334 face-to-face interviews with working professionals in the Kathmandu into a decision-grade view of where artificial intelligence has actually landed in Nepal's economy — sector by sector, and layer by layer.

WHY WE RAN THIS STUDY

Nepal is in the middle of an unusual technology moment. Generative AI arrived in the country not through enterprise procurement, corporate IT budgets or government programmes, but through mobile phones — as a free consumer product, spread by social media, adopted individually and almost entirely unsupervised.

That is the opposite of how every previous wave of enterprise technology reached Nepal. It produces a market where the usual signals are misleading: adoption looks high, but capability, workplace integration and governance may not have followed. Policymakers designing a national AI framework, and businesses deciding where to invest, both need to know which of those layers are real.

Sankhya AI commissioned this survey to measure them separately rather than assume they move together.

WHAT IS NEW HERE

- **Layered measurement.** We measure awareness, understanding, personal use, use intensity, workplace deployment and workforce preparedness as six distinct things — not one.
- **A comparable sector index.** The Sankhya AI Readiness Index scores each of eight sectors on the same four pillars, allowing a like-for-like league table.
- **Governance as a measured variable.** Regulatory appetite, ethical-harm awareness and knowledge of Nepal's own AI policy draft are measured alongside adoption.

A NOTE ON PRECISION

Every figure in this report is calculated from the raw response file and carries an explicit base. Because several questions were routed or partially completed, bases vary by question between n=280 and n=334. Where a base falls below 30 — principally Manufacturing (n=9) — findings are marked as indicative and should not be read as representative of the sector.

HOW TO READ THE EXHIBITS

Every exhibit carries a number, an action title stating what it shows, a subtitle giving the question and base, and a note recording any qualification. A full list of the exhibits, with their bases, appears on page 7.

FRONT MATTER

CHAPTERS

APPENDICES AND SOURCES

LIST OF EXHIBITS

31 exhibits, every one with a stated base

Every exhibit in this report is calculated directly from the response file. Bases vary by question and are printed with each exhibit as well as listed here.

THREE BASES RECUR

n=334 — full valid sample: demographics, awareness, understanding, tool recognition.
n=295 — perception, trust and policy module.
n=280 — industry-specific module.

MULTIPLE-RESPONSE EXHIBITS

Exhibits drawn from "select all that apply" questions are marked in their notes. Their percentages sum to more than 100 by design and should be read as the share of respondents naming each item, not as shares of a whole.

SOURCE DATA

All exhibits are reproducible from the anonymised dataset and analysis code, which are available on request. See the data availability statement in Appendix C.

ABBREVIATIONS AND GLOSSARY

Terms used in this report

Definitions are operational — they describe how each term is used and measured in this study, which is not always how it is used in general discussion.

Abbreviations

AI	Artificial intelligence
API	Application programming interface
KMC	Kathmandu Metropolitan City
L&D	Learning and development
ML	Machine learning
MoCIT	Ministry of Communication and Information Technology, Government of Nepal
n	Number of respondents answering a given question (the base)
N	Total valid sample (N=334)
p	Probability value from a statistical test
pp	Percentage points
R1–R9	Recommendations 1 to 9, Chapter 9
UGC	University Grants Commission, Nepal
UN	United Nations
UNESCO	United Nations Educational, Scientific and Cultural Organization
χ^2	Chi-square test statistic

Glossary

Adoption
Whether a respondent reports using AI personally at all. A binary measure, and deliberately distinguished throughout this report from intensity, capability and workplace deployment.

Daily user
A respondent who reports using AI once a day or multiple times a day.

Deep understanding
The top category of the self-assessed understanding question: "I understand AI well and use it regularly." Self-reported and not independently tested.

Governance gap
The 45-percentage-point difference between the share of respondents who want AI regulated (91.5%) and the share aware that the Government of Nepal has published a draft national AI policy (46.4%).

Net trust
The combined share of respondents who trust AI for important decisions either completely or with human oversight.

Purposive sample
A sample selected to achieve coverage of defined groups rather than by random selection. It supports description and comparison between those groups but not projection to population totals.

Sankhya AI Readiness Index
A composite score from 0 to 100, the unweighted mean of four equally weighted pillars: literacy, personal adoption, workplace deployment and workforce preparedness. Construction is set out in Chapter 4 and Appendix A.

Sector module
The seven industry-specific questions asked only of respondents in a given sector, routed automatically from their stated occupation. Base n=280.

Statistically significant
Used in this report to mean a chi-square test of independence returning $p<0.05$. All tests referenced in the text are listed in Appendix A.

Weekly-or-more user
A respondent who reports using AI a few times a week, once a day or multiple times a day.

Workplace deployment
Frequency of AI use for sector-specific work tasks, measured in the industry module and distinct from personal adoption.

01

CHAPTER ONE

Executive summary

Kathmandu's professionals have adopted artificial intelligence at rates that would be high in any market. What they have not built is the understanding, the workplace structure or the governance to hold it.

CHAPTER 01 · EXECUTIVE SUMMARY

Adoption has outrun capability, and capability has outrun governance

Awareness of AI among the professionals we interviewed is effectively universal and personal adoption is high. But the depth of that adoption is thin and its distribution is sharply uneven. Respondents also report almost no contact with employer training or with national AI policy — the two channels through which a technology is normally governed.

98%
have heard of AI — the awareness battle is over

81%
of all respondents use AI personally; 57% weekly or more

22%
combine deep understanding with daily use

46%
know Nepal has a draft national AI policy

1. Awareness is complete; competence is not

Ninety-eight per cent of respondents have heard of AI, and 81% report using it personally. But only 39% say they understand AI well enough to use it regularly, and just 22% combine that understanding with daily use. The funnel loses more than three-quarters of its volume between "has heard of it" and "uses it capably" — **a 76-point collapse across five steps**.

2. AI arrived through the phone, not the workplace

Eighty-seven per cent first learned about AI through social media; only 8.5% through workplace or professional training and 10.7% through formal education. On the diffusion of the most consequential general-purpose technology of the decade, Nepal's institutions reached fewer than one respondent in ten. We read the rest of this report's findings — shallow skills, low policy awareness, uneven quality of use — as consequences of that channel.

3. The toolset is one product deep

ChatGPT is known to 91% of respondents. The second-best-known tool, Google Gemini, reaches 56%; no other tool exceeds 26%. Awareness of specialist and productivity tools is negligible. The practical consequence is that "AI capability" in Kathmandu today means *general-purpose chatbot literacy* — a narrow foundation on which to build sector-specific value.

4. There is a generational cliff, not a gradient

Among respondents who answered the frequency question, daily use runs at 52% in the 20–29 band and falls to 24% among those aged 30–39 and 40–49. Education level, once age is controlled for, explains almost nothing. The sample therefore splits by age band rather than by qualification. Whether that reflects ageing or generational replacement, a single survey cannot say.

5. Trust runs ahead of understanding

Fifty-eight per cent trust AI for important decisions either completely (17%) or with human oversight (41%). Yet only 39% claim to understand it well. Trust is highest precisely where scrutiny should be greatest: 86% of media and journalism respondents express net trust, against 37% in healthcare — the sector where AI error is most consequential.

6. The workforce is anxious but not acting

Sixty-one per cent are concerned that AI will affect job security in their sector, yet only 40% are actively learning new skills; 51% say they "plan to start soon". Concern converts to action at a rate of just 49%. Healthcare is the clearest example: 63% concerned, 20% acting.

7. The governance gap is the headline risk

Ninety-two per cent want AI regulated, 78% say government should lead, and 83% have already heard of AI being used for crime or unethical activity. But only 46% know Nepal has published a draft national AI policy — **a 45-point gap between regulatory demand and awareness of the regulatory response**. The mandate for AI governance is close to unanimous; public knowledge of the response is not. Note that this study measured awareness of national policy, not whether employers have AI policies of their own — a question the instrument did not ask.

THE ONE NUMBER TO REMEMBER

22%. That is the share of respondents who both report understanding AI well and use it daily. Everything else in this report — the sector gaps, the trust anomalies, the low policy awareness — sits in the distance between that number and the 98% who have heard of AI. Both figures are self-reported; neither was independently tested.

CHAPTER 01 · IMPLICATIONS

What this means, and what to do about it

The findings point to three structural interventions — one for government, one for employers, one for the AI industry in Nepal. None of them are about raising awareness. Awareness is finished.

EXHIBIT 1

The adoption funnel loses 76 points between hearing about AI and using it capably

Share of all valid respondents (n=334) reaching each stage of the AI adoption ladder

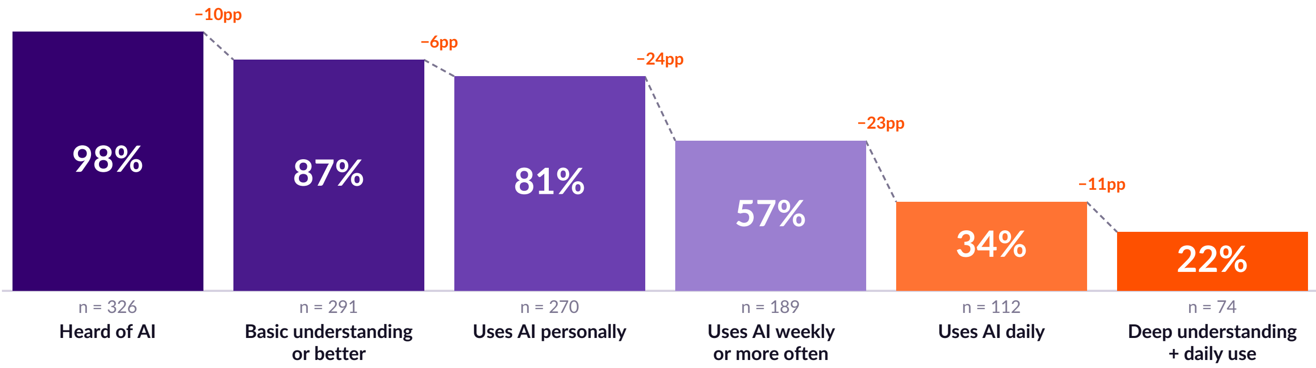

Base: all valid respondents, n=334. Percentages calculated on the full sample so that stages are directly comparable; question-level bases vary between n=289 and n=334 because of routing and item non-response. "Deep understanding + daily use" combines the top understanding category with use of AI once a day or more.

FOR POLICYMAKERS

Publish, translate and campaign — not just draft

Nepal has a draft AI policy and a population that overwhelmingly wants one. It does not have a public that knows the draft exists. Fifty-four per cent are unaware or unsure. The immediate, low-cost intervention is a distribution problem, not a drafting problem: Nepali-language summaries, sector guidance notes for healthcare and legal practice, and dissemination through the same social channels that carried AI into the country in the first place.

Sequence matters. Regulation that arrives before public understanding of what is being regulated tends to be either ignored or resisted.

FOR EMPLOYERS

Convert shadow use into supervised use

AI is already inside Nepali organisations: 94% of respondents who answered the workplace module use AI for work tasks at least occasionally, while only 8.5% learned about AI at work. Most workplace AI use in this sample is therefore self-taught. Whether it is also ungoverned we cannot say — the instrument did not ask about employer policy — but the risk is concentrated where the tasks are sensitive: 75% of banking respondents name cybersecurity as their principal concern, and 53% of healthcare respondents use AI to handle patient records.

The first move is not a training budget. It is an acceptable-use policy and a register of what staff are already doing.

FOR THE AI INDUSTRY IN NEPAL

Sell depth, not access

The market does not need to be told AI exists. It needs to be moved from general-purpose chatbot use to sector-specific application. The commercial opportunity sits in the 44% "practical adopters" — professionals who already use AI weekly, sit in the 30–39 age band, hold senior-enough positions to buy, and are not yet fluent.

The clearest underserved demand is in Legal (index 48.5, the lowest of eight sectors, but 93% expecting AI to change their industry) and Healthcare (index 57.1, with only 20% actively upskilling against 63% expressing job-security concern).

CHAPTER 01 · SCOPE OF THE EVIDENCE

What this study can say, and what it cannot

A survey of this design supports some kinds of statement and not others. This page states the boundary explicitly so that it does not have to be re-litigated at every exhibit.

WHAT THIS STUDY CAN SAY

- **What these 334 professionals report doing.** Awareness, self-assessed understanding, personal use and frequency, workplace use, upskilling behaviour, trust, concerns and knowledge of national AI policy — each on a stated base.
- **How those things differ between groups in this sample.** Differences by age, sector, gender, education and organisation size, with effect sizes and corrections for multiple testing reported in Appendix A.
- **Where the gaps are widest,** and therefore where an intervention would have the most room to work.

WHAT THIS STUDY CANNOT SAY

- **Whether AI has delivered any benefit.** The instrument asked what people use, never what they gained. No figure here is a productivity estimate.
- **What is true of Nepal, or of the Kathmandu Valley.** This is a purposive sample of 334 professionals in Kathmandu district. It supports no population projection.
- **Whether people are as capable as they say.** Understanding, trust and preparedness are self-reported and untested.
- **Whether employers have AI policies.** Not asked. Statements about workplace governance in this report concern what respondents learned and did, not what their organisations require.
- **Anything about causation.** Every relationship reported is an association measured at one moment in time.

"Nepal did not adopt AI. Nepal's twenty-somethings did, on their phones, and brought it to work with them."

Sankhya AI analysis of 334 professional interviews, Kathmandu, March–April 2025 · An interpretation of the age pattern in Exhibit 10, not a measured finding

02

CHAPTER TWO

Approach and sample

How the study was designed, who was interviewed, and what we did to the raw data before analysing it.

In this chapter · Research design · Sample composition · Sector coverage · Data cleaning decisions · What the sample can and cannot support

CHAPTER 02 · APPROACH

A structured, sector-routed instrument delivered face to face

The survey was designed to separate the layers of AI engagement that are usually collapsed into a single "adoption" number — and to ask every respondent the same general questions before routing them into a module specific to their industry.

EXHIBIT 2

The sample skews young, well-educated and early-career — as Kathmandu's professional workforce does

Composition of valid sample on four demographic dimensions, % of n=334

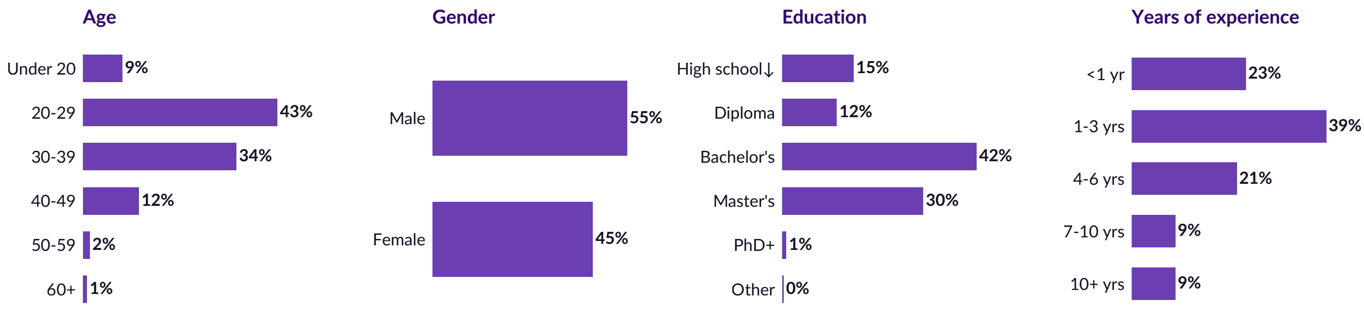

Base: n=334. Percentages may not sum to 100 because of rounding. 72.8% of the sample holds a Bachelor's degree or higher; 61.7% has three years of work experience or less; 75.4% works in an organisation of 50 people or fewer.

Research design

Sankhya AI conducted 343 structured face-to-face interviews in Kathmandu district between March and April 2025, using a team of three trained field enumerators working from a bilingual (Nepali and English) instrument. The questionnaire ran to 24 general questions plus a sector-specific module of seven further questions, routed automatically on the basis of the respondent's stated industry.

The instrument was organised in five sections: demographic profile; general awareness of AI; perception, trust and concerns; AI policy and regulation; and the industry-specific module. This structure allows the report to distinguish between *knowing about* AI, *understanding* it, *using* it personally, *using it at work*, and being *prepared* for its consequences — five things that are routinely reported as one.

Sampling approach

Respondents were selected purposively to achieve coverage across eight target sectors, with quotas designed to secure a workable base in each. Recruitment took place at workplaces and institutions across Kathmandu district, which accounts for all 334 valid respondents. Within it, 94.0% of workplace locations resolve to Kathmandu Metropolitan City and the remainder to adjoining local bodies.

This is a purposive sample, not a probability sample of any defined population, and confidence intervals are therefore not reported. Associations highlighted in the text are supported by the tests set out in Appendix A.

One composition point matters throughout. The report calls its respondents working professionals, and 90.1% are. But 33 of the 56 Education respondents — 9.9% of the sample — are students. Where the report discusses Education it is describing a mixed population of learners and educators, and says so at each point.

SURVEY AT A GLANCE

Gross contacts	**343**
Records removed (aborted)	9
Valid analysis base	**334**
Completion rate	97.4%
Field enumerators	3
Sectors covered	8
Questions (general)	24
Questions (sector module)	7
Fieldwork	Mar–Apr 2025
Geography	Kathmandu

INSTRUMENT STRUCTURE

Sections 1–2. Profile and awareness: demographics, sector, tenure, organisation size; awareness of AI, how it was first encountered, depth of understanding, tool recognition, frequency of use.

Sections 3–4. Perception and policy: benefits, trust, concerns, awareness of unethical use, preferred degree of regulation, who should regulate, awareness of Nepal's draft AI policy.

Section 5. Industry module, routed on sector: role, workplace usage, tasks, expected impact, job-security concern, upskilling, sector-specific concerns.

CHAPTER 02 · COVERAGE AND DATA TREATMENT

What the data can support, and what was done to it first

The raw export required five substantive corrections before analysis. All of them are documented, reversible and reported here rather than in a footnote.

EXHIBIT 3

Sector coverage

Share of valid sample, n=334

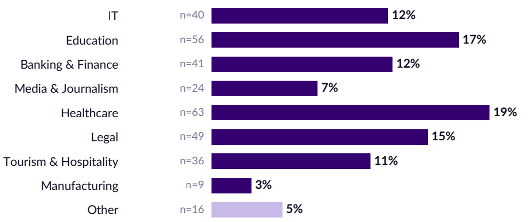

"Other" comprises respondents whose occupation fell outside the eight target sectors; they are included in valley-wide figures but excluded from sector comparisons.

Data treatment

Nine records were removed because they contained no substantive responses beyond the enumerator identifier — aborted or non-consenting contacts. The analysis base is therefore **n=334 valid records from 343 gross contacts**, a completion rate of 97.4%.

Four further corrections were applied to the raw export. Two ordinal fields — years of experience and organisation size — had been corrupted by spreadsheet date auto-formatting ("1–3 years" stored as "01-Mar"; "10–50" stored as a 1950 date serial). These were restored to their questionnaire categories.

Free-text municipality entries, which contained thirty-four distinct spellings, were normalised to local-body level, with uninterpretable entries left unclassified rather than defaulted into the largest category. Trailing whitespace in the policy-awareness field, which had split 137 affirmative responses into a separate category, was stripped. Response labels carried machine-translation artefacts from the Nepali instrument and were restored to the English questionnaire wording. All corrections are documented in Appendix A.

WHAT THE SAMPLE SUPPORTS — AND WHAT IT DOES NOT

The bases in Education (n=56), Healthcare (n=63), Legal (n=49), Banking & Finance (n=41), IT (n=40), Tourism & Hospitality (n=36) and Media & Journalism (n=24) are sufficient to support sector-level description and cross-sector comparison. **Manufacturing (n=9) is not.** It is retained in the index for completeness and marked throughout as indicative only. Readers should also note the sample's youth: 51.5% are under 30. Because age is the single strongest predictor of AI fluency in this data, valley-wide adoption figures should be treated as an upper bound on the wider working population.

BASE SIZES VARY BY QUESTION

Three bases recur through this report. **n=334** is the full valid sample and applies to demographics, awareness, understanding and tool recognition. **n=295** applies to the perception, trust and policy module, which 39 respondents did not reach. **n=280** applies to the industry-specific module.

WHY THE BASES DIFFER

The instrument was administered sequentially and some interviews ended early. Rather than impute or drop those respondents entirely, each figure is reported on the base that answered its own question, and that base is stated alongside every exhibit.

ONE VARIABLE EXCLUDED

The question on what types of task respondents use AI for was answered by only 64 respondents, apparently because of a routing fault. It is excluded from this report; the sector-specific task questions cover the same ground on a sound base.

03

CHAPTER THREE

The adoption funnel

Awareness is universal. Understanding is partial. Habitual, capable use is rare. The distance between those three states is the central fact of Nepal's AI market.

CHAPTER 03 · AWARENESS AND UNDERSTANDING

Everyone has heard of AI. Fewer than two in five say they understand it

Awareness has reached saturation: 98% of respondents have heard of artificial intelligence and only 1.8% have not. But awareness is the cheapest layer of the funnel, and in Kathmandu it has been achieved without any of the layers that normally accompany it.

EXHIBIT 4

Half the market sits in the "basic understanding, low use" band

"How well do you understand what AI can do?" — % of n=334

Base: all valid respondents, n=334. Question wording follows the English instrument; response options were translated from the Nepali field version.

The single largest group in the sample — 48.5%, or 162 people — describes itself as having "a basic understanding, but doesn't use much". This is the defining segment of Nepal's AI market: aware, unintimidated, occasionally curious, and not yet converted into a capable user. A further 12.9% sit below it, having heard of AI without understanding it (9.6%) or having no idea what it does (3.3%).

That leaves 38.6% — 129 people — claiming deep understanding and regular use. This is a substantial figure by any international standard, and it is the reason Kathmandu's AI market looks healthier from the outside than it is. The qualifier matters: this is *self-assessed* understanding, measured in a population where 87% learned about AI from social media and only 8.5% received any workplace training. There is no independent test of competence behind the claim.

The more defensible read is the combination measure. Cross-tabulating self-reported deep understanding against actual daily use gives **22.2% (n=74)** — respondents who both claim fluency and demonstrate the behaviour that usually accompanies it. That is the realistic size of Kathmandu's capable AI user base today.

98%

have heard of AI (n=326 of 334)

38.6%

claim deep understanding and regular use

22.2%

claim deep understanding *and* use AI daily

WHY THE GAP BETWEEN 38.6% AND 22.2% MATTERS COMMERCIALLY

Products, training programmes and policy interventions sized against the 38.6% figure will over-estimate the capable market by roughly three-quarters. Sized against 22.2%, they will correctly identify a small, young, concentrated user base — and correctly identify the 48.5% "basic understanding" band as the growth opportunity rather than the current market.

CHAPTER 03 · CHANNELS AND TOOLS

AI reached Nepal through the phone, and it brought one product with it

The diffusion channel explains the shape of the market. Where a technology enters through consumer social media rather than through employers or educators, it produces wide, shallow, unsupervised adoption — exactly what this data shows.

EXHIBIT 5

Institutions were absent from the diffusion of AI

"How did you first learn about AI?" — multiple responses, % of n=328

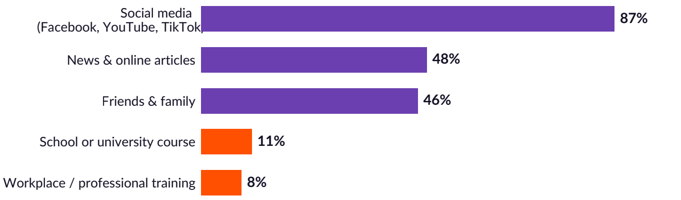

Base: respondents answering the source question, n=328. Multiple responses permitted, so percentages sum to more than 100.

The institutional vacuum

Eighty-seven per cent of respondents first encountered AI through social media — Facebook, YouTube and TikTok. News and online articles reached 48%, and friends and family 46%. Formal channels barely register: **10.7% learned about AI through school or university, and 8.5% through workplace or professional training.**

The ratio is roughly ten to one in favour of the algorithmic feed over the institution. That has three consequences that run through the rest of this report. First, what people know about AI is what performs well on social platforms, which favours spectacle over method. Second, because employers were not the channel, employers have no visibility into what their staff learned or how they are applying it. Third, because education systems were not the channel, there is no baseline curriculum to build on — the 34% of education-sector respondents actively learning AI skills are doing so on their own initiative.

EXHIBIT 6

One tool is universal; everything else is a minority

"Which of these AI tools have you heard of?" — multiple responses, % of n=334

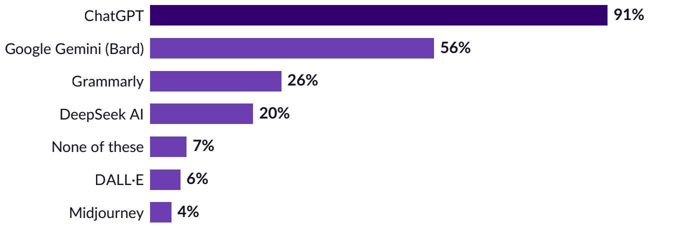

Base: all valid respondents, n=334. Multiple responses permitted. 7.2% have heard of none of the listed tools.

The one-product toolset

ChatGPT is known to 90.7% of respondents. Google Gemini follows at 56.3%, and then the distribution collapses: Grammarly 26.0%, DeepSeek AI 20.4%, DALL·E 6.0%, Midjourney 4.2%.

Two implications follow. The first is capability: awareness of image, video and specialist tools is so low that AI competence in Kathmandu effectively means competence with a general-purpose text chatbot. Sectors whose value from AI depends on specialist tooling — manufacturing quality control, medical imaging, legal document review at scale — have almost no installed base of awareness to build from.

The second is concentration risk. A market in which a single foreign product holds 91% mindshare is a market with no negotiating position on price, data residency or service continuity. DeepSeek's 20.4% recognition, achieved within months of launch, shows how quickly that position can be contested — but also how completely the market follows whatever the feed surfaces.

Ten times as many people learned about AI from a social feed as from their employer or their university.

Exhibit 5 · 87.2% social media vs 8.5% workplace training, 10.7% formal education

CHAPTER 03 · INTENSITY

Two-thirds of users are weekly or more — but a fifth barely use it at all

Frequency separates the market more cleanly than any demographic variable. It is also the layer where the sector differences that structure the rest of this report first become visible.

EXHIBIT 7

Use is bimodal: a committed daily core, and a long tail of occasional users

"How often do you use AI?" — % of n=289 who answered the frequency question

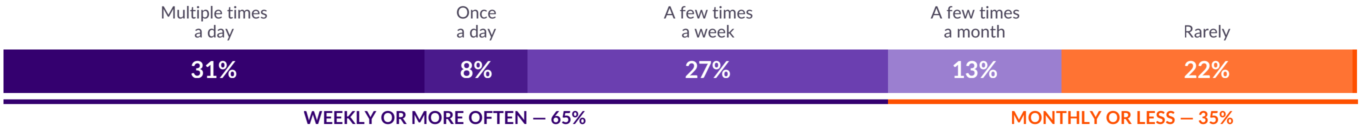

Base: n=289. Respondents who reported not using AI personally were not routed to this question, so the base is smaller than the full sample. Read against the full sample of 334, weekly-or-more users represent 56.6%. Displayed components are rounded and may not sum exactly to the bracketed totals.

Among those who use AI, 31.1% use it multiple times a day and a further 7.6% once a day — a committed core of 38.7%. Add the 26.6% using it a few times a week and 65.4% of users are weekly or more. Against the full sample of 334, that is 56.6%: a majority of Kathmandu's professionals now touch AI in a typical week.

The tail is equally important. Twenty-two per cent of users describe their usage as "rarely", and 12.8% as a few times a month. These are people who have crossed the adoption threshold without forming a habit — the group most likely to be counted as "adopters" in a headline statistic and least likely to be generating any value from the technology.

Frequency also separates the sectors sharply, and does so more decisively than awareness does. Multiple-times-a-day use runs at 54.5% in Media & Journalism and 51.4% in IT, against 16.7% in Legal and 10.3% in Healthcare. The sectoral difference in frequency is statistically significant at the 0.001 level ($\chi^2=87.6$, $p<0.001$), while the difference in bare awareness is not meaningful at all — every sector is at or near saturation.

This is why the next chapter measures readiness rather than adoption. Adoption is a binary that everyone now passes. Readiness is a composite that separates them.

WHERE DAILY USE CONCENTRATES

SECTOR	MULTIPLE TIMES/DAY
Media & Journalism	**54.5%**
IT	**51.4%**
Banking & Finance	44.7%
Education	31.4%
Tourism & Hospitality	28.6%
Legal	16.7%
Healthcare	10.3%
Manufacturing (n=9)	0.0%

Base: n=276 answering both sector and frequency. $\chi^2=87.6$, $p<0.001$.

04

CHAPTER FOUR

The Sankhya AI Readiness Index

One composite score, four pillars, eight sectors — built so that a health worker and a software developer can be compared on the same terms.

In this chapter · How the index is constructed · The league table · Pillar-level diagnosis · What separates the leaders from the laggards

CHAPTER 04 · CONSTRUCTION

Four pillars, equally weighted, scored 0–100

Adoption alone is a poor guide to readiness: a lawyer who consults a chatbot once a month and a developer who ships code with one are both "adopters". The index scores each sector on four independent dimensions and averages them.

Why a composite

Single-question rankings produce unstable and misleading league tables. Ranked on the yes/no personal-use question alone, IT and Manufacturing both score 100% — an artefact of a binary question and, in Manufacturing's case, a base of nine. Ranked on self-reported deep understanding alone, Manufacturing (55.6%) leads every sector in the sample, which is not credible.

The Sankhya AI Readiness Index therefore combines four separately measured pillars, each converted to a 0–100 scale on the sector's own answering base, then averaged with equal weight. Equal weighting is deliberate: there is no defensible empirical basis in this dataset for privileging one pillar over another, and equal weights make the index transparent and reproducible.

The four pillars

PILLAR	WHAT IT MEASURES
Literacy	Depth of self-assessed understanding of what AI can do, weighted from "no idea" (0) through "basic understanding" (0.5) to "deep understanding and regular use" (1.0)
Personal adoption	Intensity of personal AI use, weighted from never (0) and rarely (0.25) through monthly (0.5) and weekly (0.75) to daily (1.0)
Workplace deployment	Frequency of AI use for sector-specific work tasks, on the same intensity weighting, taken from the industry module
Workforce preparedness	Whether the respondent is actively learning AI skills (1.0), plans to start soon (0.5), or sees no need (0)

Pillar bases vary by sector and question; all bases are reported in Appendix B. Sector scores use only respondents who answered the relevant question.

Reading the score

A sector scoring 100 would consist entirely of people who understand AI deeply, use it daily both personally and at work, and are actively building new skills. A sector scoring 0 would have none of those things. Real sectors fall between: the valley-wide average is **63.9**.

VALLEY-WIDE PILLAR SCORES

PILLAR	SCORE	BASE
Literacy	**64.5**	318
Personal adoption	**66.0**	293
Workplace deployment	**59.6**	280
Workforce preparedness	**65.4**	280
Composite index	**63.9**	–

The valley profile is revealing in itself. Personal adoption (66.0) sits comfortably above workplace deployment (59.6): Nepal's professionals are further ahead as individuals than as employees. Preparedness (65.4) is buoyed by intention rather than action — 51.4% of the preparedness base say they "plan to start soon" against 39.6% actually learning, and the scoring treats intention as half-credit.

HOW TO USE THE INDEX

The index is a diagnostic, not a scoreboard. A high score does not mean a sector is deriving value from AI; it means the raw conditions for value are present. A low score identifies which specific pillar is the binding constraint — and therefore which intervention will move the sector fastest. Legal's problem is adoption (41.7); Healthcare's is workplace deployment (47.5) and preparedness (53.4). Those require entirely different responses.

CHAPTER 04 · THE LEAGUE TABLE

A 31-point spread separates Nepal's most and least AI-ready sectors

IT and Media & Journalism lead by a clear margin. Legal sits 31 points below IT and 15 points below the valley average — an outlier that the sector deep dive examines in detail.

EXHIBIT 8

Sankhya AI Readiness Index — eight sectors ranked

Composite score, 0–100, average of four equally weighted pillars; sector bases n=9 to n=63

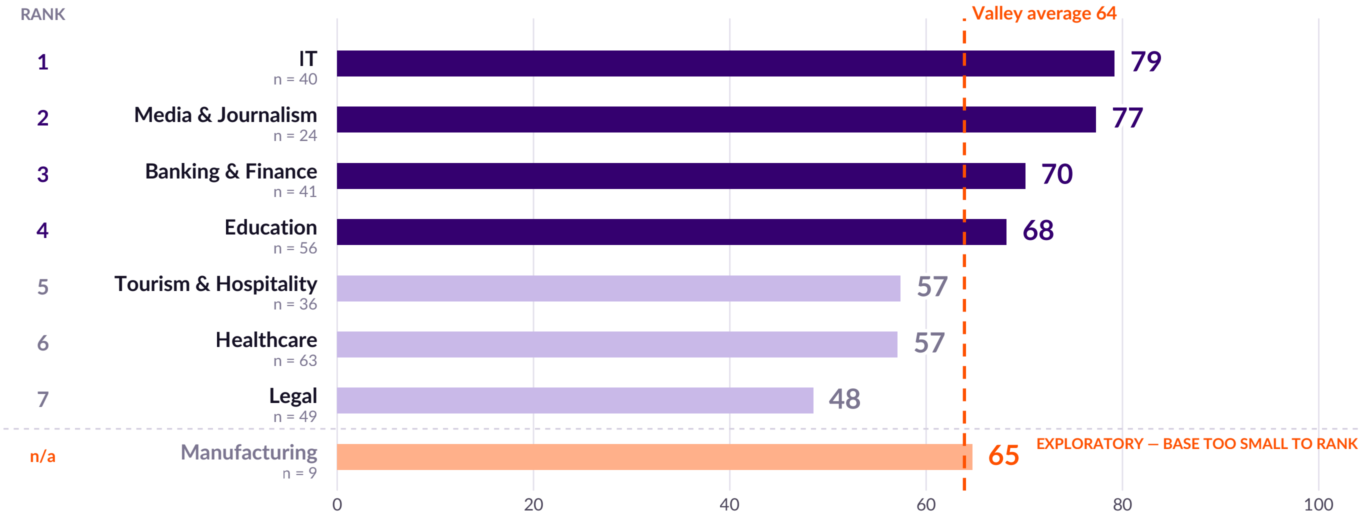

Sector bases as shown. Manufacturing (n=9) is indicative only and should not be treated as representative. Valley average (63.9) is calculated across all respondents in the eight target sectors, not as a mean of the eight sector scores.

READING THE TABLE

The eight sectors divide into three bands. IT and Media & Journalism form a clear leading group above 77. Banking & Finance, Education and Manufacturing sit in a middle band between 64 and 71, at or just above the valley average of 63.9. Tourism & Hospitality, Healthcare and Legal trail between 48.5 and 57.4.

WHAT THE RANKING IS NOT

The index measures conditions, not returns. No sector in this study was asked what AI has actually delivered — hours saved, errors avoided, revenue gained — and none of these scores should be read as a productivity claim. A sector can score highly on readiness and still be extracting little value, which the application patterns in Chapter 7 suggest is the case in several.

WHERE THE GAPS ARE WIDEST

Among sectors with a workable base, the widest internal pillar spread is IT's — 22.5 points between preparedness (88.6) and literacy (66.1) — followed by Healthcare's 20.1 points between literacy (67.6) and workplace deployment (47.5). The widest spread between sectors on a single pillar is workplace deployment: 43.4 points between Media & Journalism (81.8) and Legal (38.4).

CHAPTER 04 · PILLAR DIAGNOSIS

Every sector has a different binding constraint

The composite score says which sectors are ahead. The pillar breakdown says why — and therefore what would move each one.

EXHIBIT 9

Pillar diagnosis: where each sector is strong and where it is constrained

Pillar score by sector, 0–100, ordered by composite index; pillar bases n=8 to n=63

	Literacy	Personal adoption	Workplace deployment	Workforce preparedness
IT	66	84	79	89
Media & Journalism	71	80	82	77
Banking & Finance	71	76	66	68
Education	67	76	66	64
Tourism & Hospitality	51	58	59	61
Healthcare	68	60	48	53
Legal	57	42	38	57
Manufacturing	68	66	50	75

Darker = stronger. Manufacturing's workplace and preparedness scores rest on n=8 and should not be compared with the other sectors.

What separates the leaders

IT (79.2) and Media & Journalism (77.3) lead not because they are more aware of AI — every sector is near saturation on awareness — but because they are the only two sectors where all four pillars clear 66. IT's preparedness score of 88.6 is the highest single pillar reading anywhere in the study: 80% of IT respondents are actively learning new AI skills, against a valley average of 39.6%.

Media & Journalism's position is more precarious than its rank suggests. Its workplace deployment score (81.8) is the highest in the sample, and 63.6% use AI daily in their work. But it also records the highest job-security anxiety of any sector — 36.4% "very concerned", more than double the next highest. This is a sector adopting quickly under duress.

What holds the laggards back

Legal (48.5) is the clearest case of a single binding constraint. Its literacy score of 56.9 is low but not disastrous; its adoption score of 41.7 is the lowest in the study by 16 points. Legal professionals know about AI and expect it to change their industry — 93.0% expect significant change or meaningful improvement — but 62.8% describe their workplace use as "occasional" and only 11.6% use it daily. Adoption, not awareness, is the gap.

Healthcare (57.1) has the opposite profile: literacy of 67.6 — above the valley average of 64.5, and above Tourism and Legal — but workplace deployment of 47.5 and preparedness of 53.4. Healthcare workers understand AI; their institutions have not given them a way to use it, and only 20.3% are actively building skills.

Tourism & Hospitality (57.4) is the only sector where literacy (51.0) is the weakest pillar. It is a foundational knowledge problem rather than a deployment or motivation problem, and it responds to different treatment.

05

CHAPTER FIVE

Six forces shaping the market

Cross-cutting patterns that hold across every sector: the generational cliff, the intensity gap between men and women, and the widening distance between anxiety and action.

CHAPTER 05 · FORCE ONE

Age is the strongest predictor of AI fluency — and education is not

The drop in daily use between the twenties and the thirties is not a gradient. It is a cliff, and it falls exactly where organisational authority begins.

EXHIBIT 10

Daily use halves between the twenties and the thirties

% within each age band, n=334

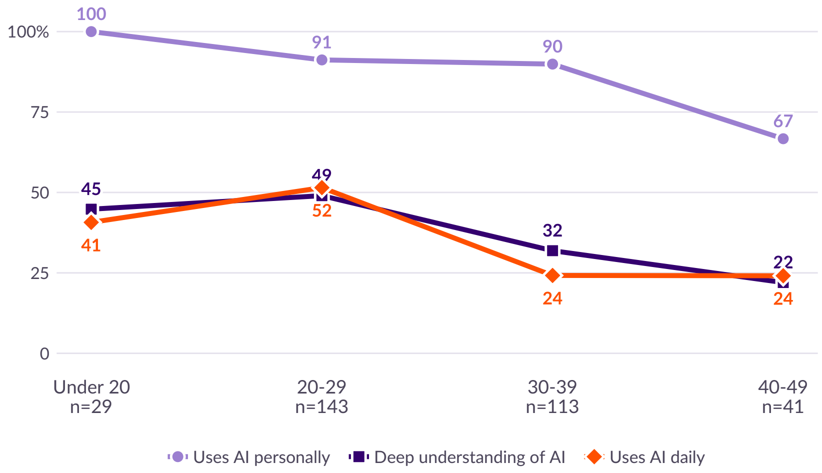

Each series is calculated on the base that answered its own question: understanding n=334; personal use n=310; frequency n=289. Sector bases as shown are the number of respondents in each age band. The 50–59 (n=5) and 60+ (n=3) bands are omitted — three respondents cannot support a trend line. Age × use frequency, collapsed to three bands: Cramér's V=0.19, Monte Carlo p<0.001, survives correction (q<0.001). Age × deep understanding: V=0.22, p<0.001, survives correction.

Personal adoption declines with age, but gently. Among those who answered the question it runs at 100% below 20, 91.2% in the twenties, 89.9% in the thirties and 66.7% in the forties. Only in the forties does it fall away sharply.

Daily use behaves differently. It runs at 51.5% in the 20–29 band and then **falls by more than half, to 24.2% among those aged 30–39**, staying at 24.1% in the forties. The step between the twenties and the thirties is far sharper than the adoption curve underneath it, and the association between age and use frequency is the second-strongest in the study (Cramér's V=0.19, Monte Carlo p<0.001, surviving correction for multiple testing).

The implication is uncomfortable for employers, but it requires one inference the survey does not itself support. This study measured age, tenure and organisation size — not seniority, budget authority or decision rights. *If* the common assumption holds that authority in Nepali organisations rises with age, then the cohort that approves budgets is also the cohort least likely to use AI habitually. We flag that as an inference rather than a finding. It is testable, and the next wave should test it directly.

EDUCATION DOES NOT EXPLAIN WHAT AGE EXPLAINS

At first reading the data appears to show a perverse education effect: 46.0% of Bachelor's degree holders report deep understanding of AI against 29.7% of Master's degree holders. That inverts once age is controlled for. Master's holders in this sample are substantially older — 49.5% are aged 30–39 and 23.8% are 40–49, against 33.1% and 6.5% respectively among Bachelor's holders. Within the under-30 group, deep understanding runs at 51.8% for Bachelor's holders and 40.9% for Master's holders; within the 30-and-over group, 37.5% and 26.6%. The residual education effect is small and the base within cells is thin.

The operative conclusion: in this sample, age is associated with AI fluency and formal qualification is not. Because the study measures a single moment in time it cannot separate an age effect from a cohort effect — whether people become less fluent as they age, or whether this particular generation entered work already fluent. Either way, screening candidates by qualification will not identify AI capability. Employers should assess the skill directly rather than infer it from a degree — or from age, which would be unreliable at the individual level and unlawful in many jurisdictions.

WHY THE CLIFF MAY FALL WHERE IT DOES

The 20–29 cohort entered work as generative AI became free and mobile-accessible; the 30–39 cohort had settled working methods before it arrived. Exposure is a more plausible explanation than aptitude — though the survey tests neither.

WHY IT MAY MATTER MORE IN NEPAL

Where employers drive adoption, seniority and fluency tend to move together because training is delivered top-down. Where adoption is consumer-led, they need not. Nepal is in the second category — though, as noted, this study did not measure seniority.

WHAT IT IMPLIES FOR TRAINING

Open-enrolment training will be taken up disproportionately by the cohort that needs it least. Reaching the 30–49 group calls for targeted, role-specific programmes with protected time.

CHAPTER 05 · FORCE TWO

Men and women are equally aware of AI. They use it very differently

There is no gender gap in awareness or in whether people use AI at all. There is a substantial one in how intensively they use it — and it widens at every step up the frequency ladder.

EXHIBIT 11

The gender gap opens as intensity rises

% within gender; men n=183, women n=151

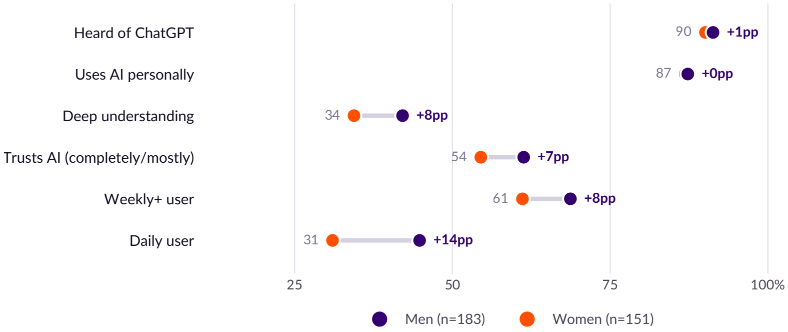

Base: n=334. Gender × personal use: χ^2=0.0, p=1.00 (no difference). Gender × use frequency: χ^2=12.3, p=0.031 (significant).

At the shallow end of the funnel there is no gender gap at all. Awareness of ChatGPT differs by 1.2 points and, among those who answered the question, personal adoption differs by 0.4 points — 87.3% of men against 86.9% of women (χ^2=0.00, p=1.00). On whether someone uses AI, men and women in this sample are indistinguishable.

The gap opens with intensity. Among respondents who answered the frequency question, weekly-or-more use runs at 68.7% for men against 61.1% for women — a 7.6-point gap. Daily use runs at 44.8% against 31.0%: **a 13.8-point gap, with men 1.4 times as likely as women to be daily users.**

That difference is present in the sample but statistically weak. Collapsed into three frequency bands so that no cell is sparse, it returns Cramér's V=0.14 with a Monte Carlo p-value of 0.059, and it does not survive correction for the sixteen tests reported in this study (q=0.079). **We report it as a descriptive pattern, not an established difference.** A larger sample would be needed to settle it.

Self-reported deep understanding follows the same shape (42.1% against 34.4%, a 7.7-point gap, V=0.07), as does trust in AI for important decisions (61.3% against 54.5% among those who answered). Neither reaches significance.

A gap that appears only at the intensity end of the funnel is unlikely to be an access or awareness problem. It is more consistent with differences in role, in discretionary time, or in confidence — and it points to intervention design that targets depth of use among women already using AI, rather than another awareness campaign.

CHAPTER 05 · FORCE THREE

The workforce is anxious about AI and largely not acting on it

Sixty-one per cent expect AI to threaten job security in their sector. Forty per cent are doing something about it. The distance between those two numbers is the clearest workforce-policy signal in this study.

EXHIBIT 12

Four sectors sit in the danger quadrant: high anxiety, low action

Job-security concern vs active upskilling, by sector; bubble size = sector base; n=280

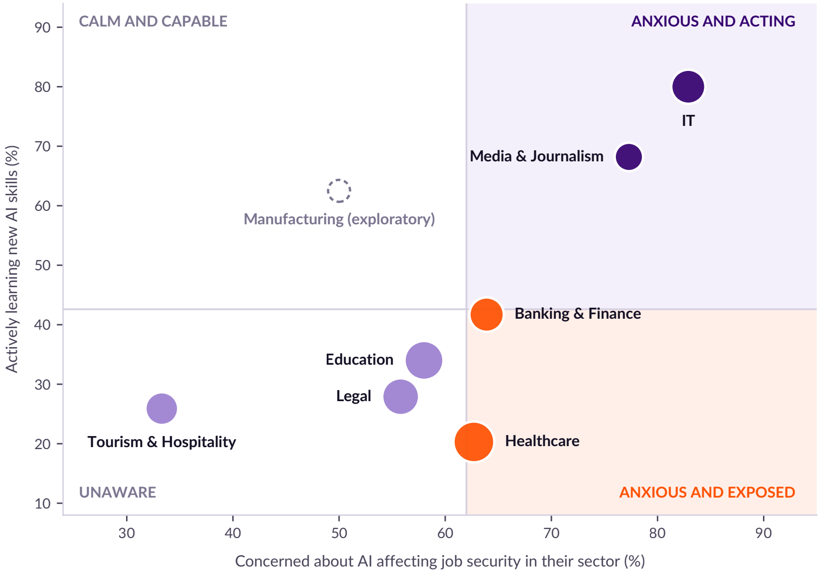

Base: n=280 answering the sector module. Quadrant boundaries are the unweighted means of the eight sector values. Sector × job-security concern: χ^2=52.6, p=0.003. Sector × upskilling: χ^2=54.5, p<0.001. Manufacturing rests on n=8.

Across the eight sectors, 61.4% are very or somewhat concerned that AI will affect job security in their industry. Only 39.6% are actively learning new skills in response; 51.4% say they "plan to start soon" and 8.9% see no need.

Concern does convert to action, but weakly. Among the concerned, 48.8% are actively learning; among the unconcerned, 25.0%. Anxiety roughly doubles the probability of upskilling (χ^2=45.1, p<0.001) — yet still leaves a majority of worried professionals doing nothing.

THE EXPOSED QUADRANT

Healthcare (62.7% concerned, 20.3% acting) and **Banking & Finance** (63.9% concerned, 41.7% acting) combine above-average anxiety with below-average preparation. Between them they account for 104 of the 334 respondents in this study — **31% of the sample sitting in the quadrant where disruption expectation exceeds disruption readiness.**

The contrast with IT is instructive. IT records the highest concern of any sector (82.9% very or somewhat concerned) and the highest action rate (80.0% actively learning). The sector closest to the technology is both the most worried about it and the most prepared. That is what a functioning response looks like.

Tourism & Hospitality sits in the opposite corner: 33.3% concerned, 25.9% acting. Low anxiety here is not reassurance — it is the lowest literacy score in the study (51.0) expressing itself as complacency.

CHAPTER 05 · FORCE FOUR

Expectations of disruption are near-universal — but rarely existential

Sectors agree almost entirely on what AI will do to their industry. They differ enormously on whether they are doing anything about it. Belief is uniform; behaviour is not.

EXHIBIT 13

Every sector expects change; almost none expects replacement

Expected impact on own sector over 2–3 years, and current upskilling, % by sector; n=280

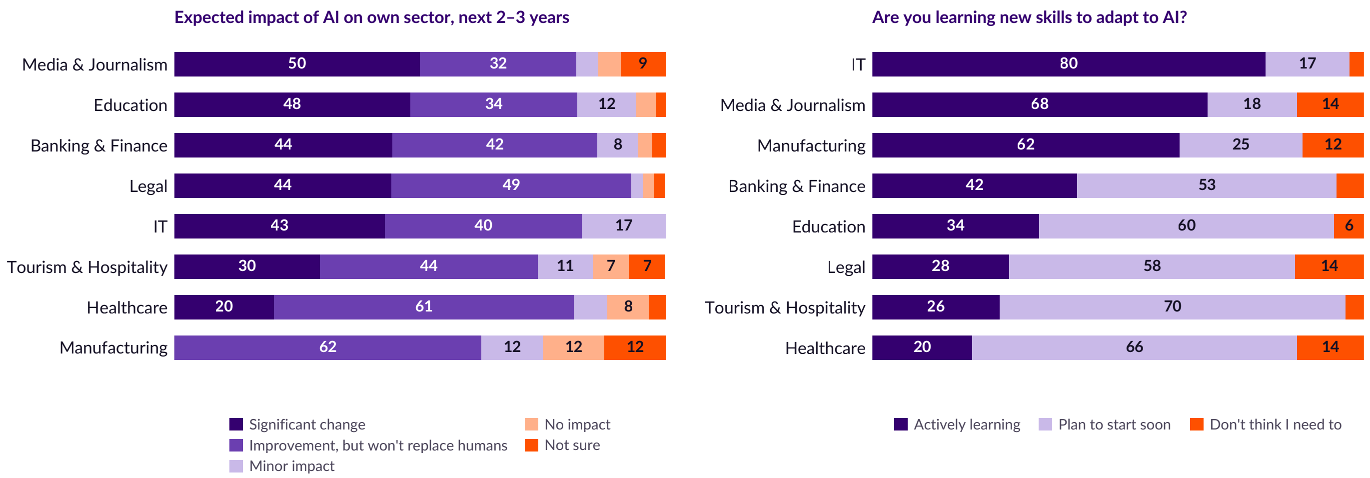

Base: n=280 answering the sector module. Sector × expected impact: χ^2=36.9, p=0.120 (not significant — expectations are broadly uniform). Sector × upskilling: χ^2=54.5, p<0.001 (highly significant — behaviour is not).

Eighty-three per cent of the sector-module base expect AI to bring either significant change (37.5%) or meaningful improvement without replacing human roles (45.4%). Only 13.5% expect minor or no impact. Notably, the differences *between* sectors on this question are not statistically significant (χ^2=36.9, p=0.120): Kathmandu's professionals broadly agree that AI is coming for their industry.

What they do not agree on is whether it will replace them. Media & Journalism records the highest expectation of significant change (50.0%) and Education is close behind (48.0%), while Healthcare is the most conservative — 61.0% expect improvement without replacement, the highest such reading in the study, and only 20.3% expect significant change.

The divergence between the two panels of Exhibit 13 is the finding. Expectation of disruption is uniform across sectors and not statistically distinguishable. Behavioural response to that expectation is *highly* differentiated (χ^2=54.5, p<0.001), ranging from 80.0% actively learning in IT to 20.3% in Healthcare.

In other words: sectors are not differentiated by what they believe about AI. They are differentiated by whether belief translates into action. That is a management and institutional variable, not an attitudinal one — and it is therefore addressable.

Across the whole sector base, only 8.9% say they see no need to learn AI skills. Denial is not the constraint. Inertia is.

CHAPTER 05 · FORCE FIVE

Trust runs ahead of understanding — and is highest where stakes are lowest

Fifty-eight per cent trust AI for important decisions. Thirty-nine per cent say they understand it. The gap between those numbers is not a rounding artefact; it is a risk profile.

EXHIBIT 14

Net trust varies 2.3× across sectors — and inversely with consequence

"How much do you trust AI in making important decisions?" — % of n=295 overall; sector panel n=280

Trust in AI for important decisions — all respondents

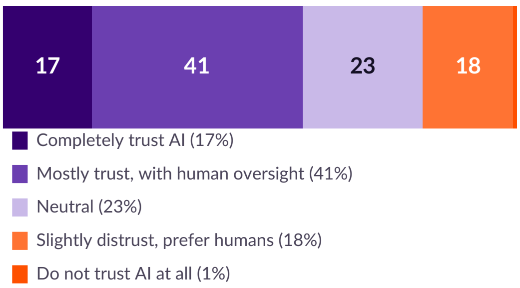

Net trust varies 2.3× across sectors

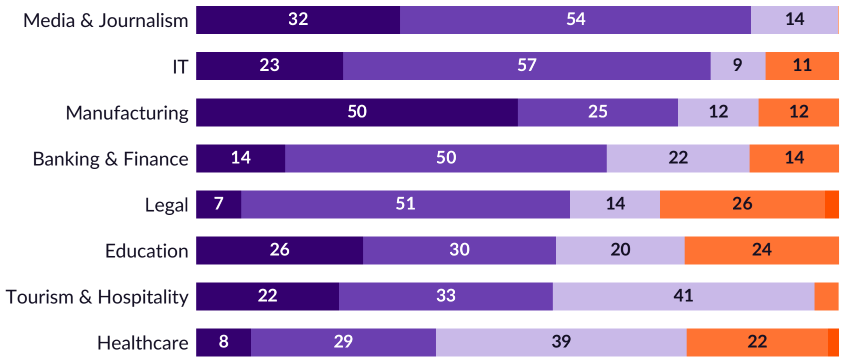

Base: n=295 answering the trust question; sector panel n=280. Sector × trust: Cramér's V=0.27, Monte Carlo p<0.001, surviving correction for multiple testing. The association between personal AI use and trust is *not* robust (V=0.09, p=0.38) and is discussed in the text.

Fifty-eight per cent of respondents extend some form of trust to AI for important decisions — 17.3% completely, 41.0% with human oversight. Only 18.3% express any distrust, and outright rejection is negligible at 0.7%. Against that, 38.6% claim to understand AI well. Trust exceeds comprehension by roughly twenty points.

The qualified-trust majority is the healthiest signal in this section. "Mostly trust, but human oversight is needed" is the correct posture toward current AI systems, and it is the single largest response at 41.0%. The concern is the 17.3% who trust AI completely — a group that, in a population with negligible formal AI training, is trusting a system it has never been taught to evaluate.

Users do appear more confident than non-users — 18.9% of AI users trust AI completely against 3.8% of non-users. But the association is not robust. Collapsed into trust, neutral and distrust bands so that no cell is sparse, it returns Cramér's V=0.09 and a Monte Carlo p-value of 0.38. **The apparent link between using AI and trusting it does not hold up in this data**, and no conclusion in this report rests on it.

THE INVERSE-CONSEQUENCE PATTERN

Net trust (completely + mostly) is **highest in Media & Journalism (86.4%)**, where AI output reaches the public directly and misinformation risk is greatest, and **lowest in Healthcare (37.3%)**, where clinical error is most consequential but human oversight is also most institutionalised.

Healthcare's caution is professionally appropriate. Media's confidence is harder to defend: the same sector reports that 54.5% see AI-generated misinformation as a principal concern in their industry, yet 86.4% extend net trust to AI for important decisions. Awareness of the risk and behavioural caution have come apart.

SECTOR	NET TRUST	DEEP UNDERSTANDING
Media & Journalism	**86.4%**	50.0%
IT	**80.0%**	47.5%
Manufacturing (n=9)	75.0%	55.6%
Banking & Finance	63.9%	43.9%
Legal	58.1%	22.4%
Education	56.0%	42.9%
Tourism & Hospitality	55.6%	25.0%
Healthcare	**37.3%**	41.3%

Legal shows the widest divergence: 58.1% net trust on 22.4% deep understanding.

CHAPTER 05 · FORCE SIX

Organisation size explains almost nothing — and that is the finding

Large employers in most markets drive technology adoption through procurement and training. In Kathmandu they do not. AI fluency is distributed almost evenly across organisations of every size, which is what you would expect of a technology that entered through the phone rather than the purchase order.

EXHIBIT 15

Deep understanding of AI barely varies with employer size

% reporting deep understanding of AI, by number of employees; n=334

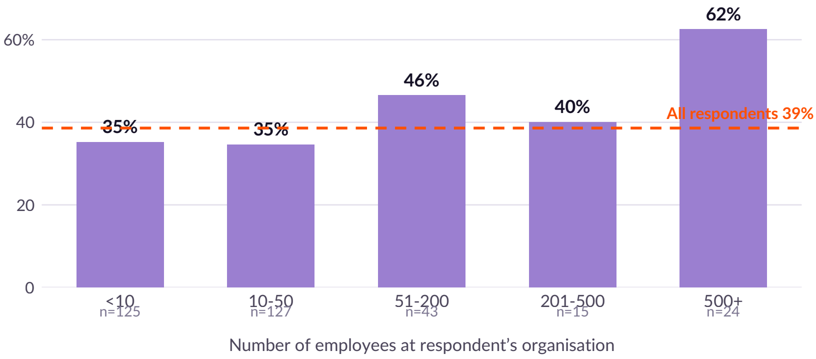

Base: all valid respondents, n=334. Organisation size × deep understanding: χ^2=8.4, Cramér's V=0.16, Monte Carlo p=0.079. The association does not reach significance and does not survive correction for multiple testing (q=0.097). The 500+ category rests on n=24.

Deep understanding runs at 35.2% in organisations of fewer than ten people and 34.6% in those of ten to fifty — statistically indistinguishable. It rises to 46.5% in organisations of 51–200 and 62.5% in those above 500, but that last figure rests on 24 respondents and the overall association is not statistically significant.

Set against the age effect in Force one — where daily use more than halves between the twenties and the thirties, with a clear and robust association — the contrast is instructive. **Who you are predicts your AI fluency in this sample. Where you work does not.**

This is the mechanism behind the institutional vacuum described in Chapter 3 seen from the employer's side. If large organisations were procuring AI, training staff and setting policy, their employees would look measurably different from those in small firms. In this sample they do not.

WHY A NULL RESULT BELONGS IN THE REPORT

A finding that nothing happens is still a finding. If AI capability in Kathmandu were being built by employers, organisation size would be among the strongest predictors in this dataset. It is among the weakest — weaker than age, sector or gender.

WHAT IT MEANS FOR PROCUREMENT

Vendors selling to large Nepali employers cannot assume an installed base of AI-literate staff. The 500-plus employer looks much like the ten-person firm. Enterprise deployment will need to carry its own training, not build on existing capability.

WHAT IT MEANS FOR POLICY

Workforce programmes that target large employers as a distribution channel will reach only a fraction of the affected workforce: 75.4% of respondents work in organisations of fifty people or fewer. Sector associations are a better route than employer size.

06

CHAPTER SIX

Risk, ethics and the governance gap

Kathmandu's professionals see the risks clearly, want regulation almost unanimously, and do not know that Nepal has already drafted a policy. That is a solvable problem — and an urgent one.

In this chapter · Benefits and risks weighed · Awareness of ethical harm · The regulatory mandate · The 45-point governance gap

CHAPTER 06 · BENEFITS AND RISKS

The market is optimistic about efficiency and worried about employment

Respondents name more benefits than concerns on average, but 93.2% name at least one concern. This is not a naïve population — it is an informed one operating without institutional support.

EXHIBIT 16

Time saving dominates the upside; jobs and privacy dominate the downside

"Biggest benefit of AI" and "major concerns about AI" — multiple responses, % of n=295

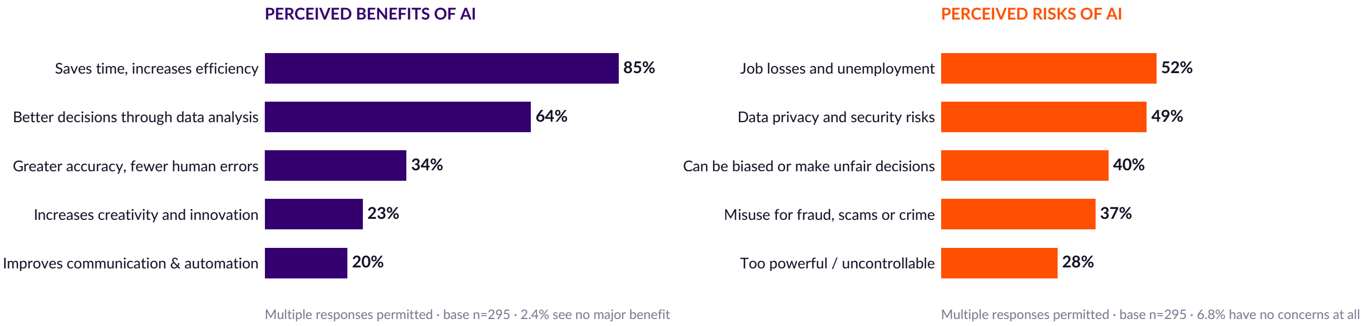

Base: n=295 answering the perception module. Multiple responses permitted. Respondents named an average of 2.25 benefits and 2.05 concerns.

The benefit case is narrow

Efficiency dominates: 84.7% name time saving as a principal benefit, and 63.7% cite data-driven decision-making. After that the case thins considerably — accuracy and error reduction at 33.9%, creativity at 23.4%, communication and automation at 19.7%.

This is a market that understands AI as a productivity tool and has not yet encountered it as a capability tool. The framing is entirely consistent with a population whose exposure is to a general-purpose chatbot: it saves time on tasks you already do. The higher-value applications — new products, new analysis, new services — are not yet in the mental model.

Only 2.4% see no major benefit at all. Scepticism is not the barrier.

The risk case is broad

Concerns are more evenly distributed, which is itself a sign of an informed population. Employment leads at 51.5%, followed by data privacy and security at 49.2%, bias and unfair decisions at 40.0%, misuse for fraud and crime at 36.9%, and loss of control at 27.8%.

Only 6.8% report no concerns whatsoever. Ninety-three per cent name at least one.

The distribution is notable for what it does *not* show. There is no single dominant fear that could be addressed with a single intervention, and no sign of the moral-panic framing that has characterised AI discourse in some markets. Kathmandu's professionals hold a differentiated risk model: economic, informational, procedural and security risks are each recognised by a substantial minority.

Ninety-three per cent name at least one concern about AI. Only forty-six per cent know that their government has drafted a policy in response.

Exhibits 16 and 18 · base n=295

CHAPTER 06 · THE GOVERNANCE GAP

An overwhelming mandate for regulation, met by an information vacuum

Nepal's professionals have already seen AI harm, already want it regulated, and already know who should do it. What they do not know is that the process has begun.

EXHIBIT 17

Ethical harm is widely recognised

"In what ways has AI been used for unethical activity?" — multiple responses, % of n=295

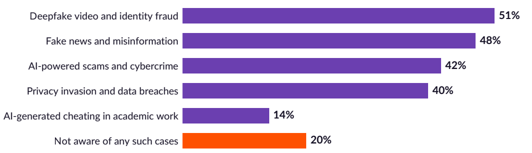

Base: n=295. Multiple responses permitted. 82.7% separately confirm having heard of AI being used for crime or unethical activity.

Deepfake video and identity fraud is the most widely recognised category of AI harm (51.2%), followed by fake news and misinformation (48.1%), AI-powered scams and cybercrime (42.4%) and privacy invasion (40.3%). Academic cheating registers at only 14.2% — strikingly low in a sample that is 16.8% education-sector and where two-thirds of the education module respondents are students.

Only 20.3% are unaware of any such cases. Separately, 82.7% confirm having heard of AI being used for crime or unethical activity. This is a population that has already encountered the harm side of the technology, largely through the same social channels that introduced them to it.

That combination — high harm awareness, low institutional contact — is what produces the regulatory pattern on the following page. The demand for governance in Nepal is not theoretical. It is a response to observed damage.

EXHIBIT 18

The governance gap: 92% want regulation, 46% know a policy has been drafted

% of n=295 answering the policy and regulation module

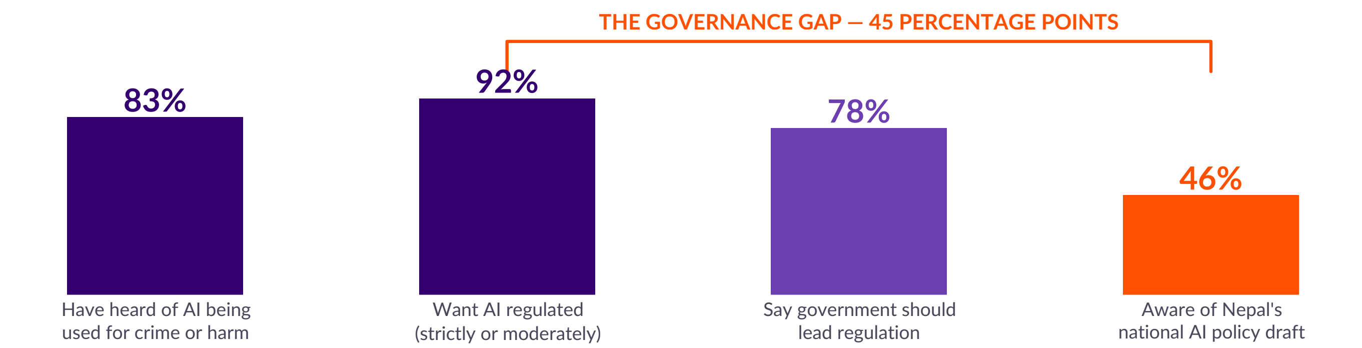

Base: n=295. "Want AI regulated" combines strict regulation (55.9%) and moderate regulation (35.6%). Only 4.1% want minimal or no regulation; 4.4% are unsure.

Fifty-five point nine per cent want AI strictly regulated to prevent misuse, and a further 35.6% want moderate regulation balancing innovation and safety. Combined, **91.5% support regulation**. Opposition is negligible: 3.4% want minimal regulation and 0.7% want none. Sector differences on this question are not statistically significant (χ^2=28.4, p=0.446) — the mandate is uniform.

Against that, 46.4% are aware that the Government of Nepal has published a draft national AI policy. Forty-one per cent say they are not, and 12.5% are unsure. **The gap between regulatory demand and awareness of the regulatory response is 45 percentage points.**

Awareness of the draft policy is also distributed counter-intuitively. It is highest in Healthcare (64.4%), Manufacturing (62.5%) and Legal (62.8%) — the three sectors with the lowest AI adoption — and *lowest* in Media & Journalism (22.7%) and Education (32.0%), the sectors with the highest use intensity (χ^2=30.2, p=0.007).

The professionals using AI most heavily are the least likely to know how their government intends to govern it. For a policy whose effectiveness depends on the behaviour of exactly that group, this is the operationally significant finding of the chapter.

CHAPTER 06 · INSTITUTIONAL EXPECTATIONS

Government first, industry a close second — and almost no appetite for self-regulation alone

Respondents assign responsibility for AI governance in a shared-accountability model, but with a clear primary actor.

EXHIBIT 19

A shared-responsibility model with the state at its centre

"Who should regulate AI?" and "What should be done to ensure ethical use?" — multiple responses, % of n=295

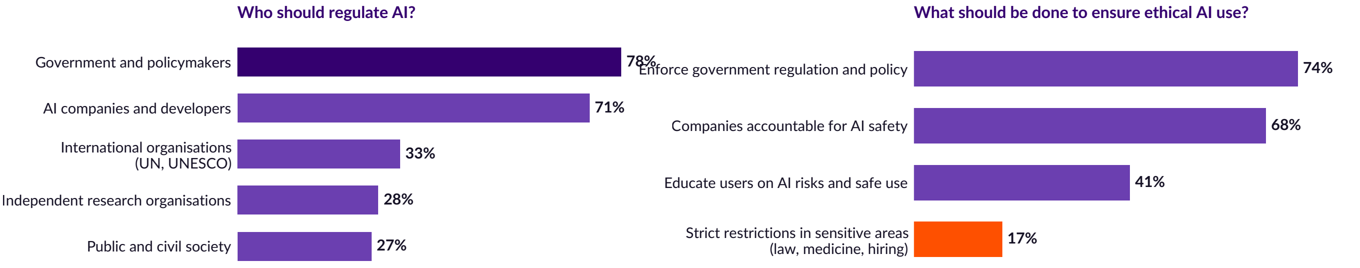

Base: n=295. Multiple responses permitted. Only 0.7% say AI should not be regulated and 1.4% say nothing should be done.

Seventy-seven point six per cent name government and policymakers as a responsible party, and 71.2% name AI companies and developers. The near-parity is significant: this is not a population that wants to hand the problem entirely to the state, nor one that trusts industry to police itself. It is asking for both.

The second tier is thinner. International organisations such as the UN and UNESCO are named by 32.9%, independent research organisations by 28.5%, and public and civil society by 27.1%. Nepal's professionals see AI governance as a national-plus-industry problem rather than a multilateral or civil-society one.

On mechanisms, 73.6% call for enforced government regulation and policy, and 67.5% for corporate accountability on AI safety. User education trails at 41.4%.

THE BLIND SPOT

Only **16.9%** support strict restrictions in sensitive domains — law, medicine, hiring. This is the lowest-supported governance mechanism in the study, and it sits directly at odds with the sector-level findings: 59.3% of healthcare respondents worry about patient data privacy, 47.5% about AI replacing medical judgement, and 69.8% of legal respondents worry about data privacy in legal cases.

Respondents recognise domain-specific risk when asked about their own domain, but do not translate that into support for domain-specific rules when asked in the abstract. Policy design should not read the 16.9% as a mandate against sectoral regulation; it more likely reflects unfamiliarity with what sectoral AI regulation would involve.

IMPLICATIONS FOR NEPAL'S AI POLICY PROCESS

1. The mandate exists. 91.5% support regulation with no meaningful sectoral variation. Political risk to acting is low.

2. Dissemination is the binding constraint. 53.5% do not know a draft exists, and awareness is lowest among heavy users.

3. Co-regulation matches expectation. 71.2% already expect developers to carry responsibility; a framework placing duties on both the state and providers will meet the public where it already is.

07

CHAPTER SEVEN

Sector deep dives

Eight one-page diagnostics, ordered by readiness index. Each combines the sector's pillar profile, its current applications, its perceived opportunity and its specific concerns.

01

SECTOR DEEP DIVE · IT

The benchmark sector — and the only one where anxiety has fully converted into action

79.2
READINESS INDEX
RANK 1 OF 7

IT tops the readiness index on the strength of adoption, deployment and preparedness rather than literacy, where three sectors score higher. Its distinguishing feature is response: it records the highest job-security concern in the study and, by a wide margin, the highest upskilling rate.

40	48%	100%	60%	80%	83%	49%
Respondents in sector	Report deep understanding of AI	Use AI personally	Use AI daily at work	Actively learning new AI skills	Concerned about job security	Aware of Nepal's AI policy draft

IT is the only sector in which all four readiness pillars exceed 66, and its workforce preparedness score of 88.6 is the highest single pillar reading anywhere in this study. Eighty per cent of IT respondents are actively learning new AI skills against a valley average of 39.6%, and just 2.9% say they see no need.

Usage is deep as well as broad. Sixty per cent use AI daily in their work — the second-highest figure after Media & Journalism — and 51.4% use AI personally multiple times a day. Personal adoption is universal: every IT respondent who answered the question reports using AI.

The application pattern is notable. The most common workplace use is data analysis and insight (62.9%), ahead of debugging and coding assistance (48.6%). IT professionals in Kathmandu are using AI as an analytical tool more than as a coding tool, which diverges from the global pattern and suggests the local market is further from AI-assisted software engineering than headline adoption implies. Only 14.3% use AI for model development or training — Nepal has adopters, not builders.

The risk profile is unusually clear-eyed. Eighty per cent name automation replacing IT jobs as a concern, and 68.6% name AI introducing security vulnerabilities. The first is the joint-highest concern reading for any single item in any sector module. This is a workforce that understands exactly what it is adopting.

SO WHAT

IT is the reference case for what a functioning sector response looks like: high concern, high action, high deployment. It is also the natural first market for advanced tooling — but the low model-development figure (14.3%) shows the ceiling. Nepal's IT sector is a sophisticated consumer of AI, not yet a producer of it.

READINESS PILLARS

PILLAR	SCORE	BASE
Literacy	**66.1**	40
Personal adoption	**83.6**	35
Workplace deployment	**78.6**	35
Workforce preparedness	**88.6**	35
Composite	**79.2**	–

WHO WAS INTERVIEWED

Software developer 26% · **Data scientist** 26% · **Other** 20% (base n=35)

EXPECTATIONS AND RESPONSE

Expect significant change	**42.9%**
Expect change of some kind	**82.9%**
Actively learning AI skills	**80.0%**
Net trust in AI	**80.0%**
Never use AI at work	2.9%

Sector module base n=35; some respondents did not complete the module.

EXHIBIT 20

IT: what AI is used for, where it could add value, and what worries the sector

Multiple responses permitted; base n=35

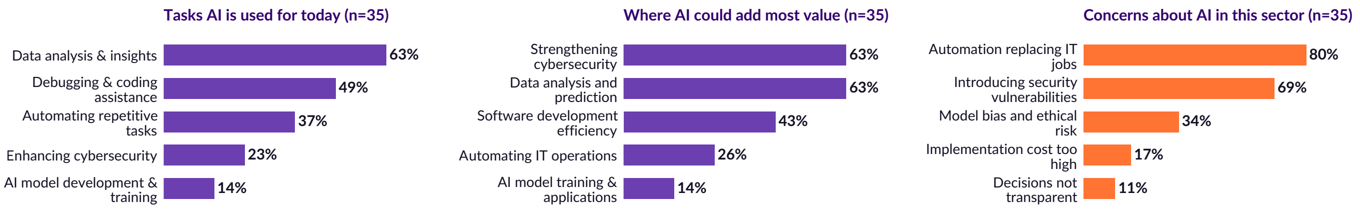

Base: IT respondents completing the sector module, n=35. Top six options shown in each panel.

02

SECTOR DEEP DIVE · MEDIA & JOURNALISM

The fastest adopter and the most anxious workforce in Nepal

77.3

READINESS INDEX
RANK 2 OF 7

Media records the highest workplace deployment score, the highest net trust and the highest rate of severe job-security anxiety in the study. It is adopting AI at speed and under duress.

24	50%	83%	64%	68%	77%	23%
Respondents in sector	Report deep understanding of AI	Use AI personally	Use AI daily at work	Actively learning new AI skills	Concerned about job security	Aware of Nepal's AI policy draft

Media & Journalism has the highest workplace deployment pillar in the study (81.8) and the highest rate of daily on-the-job AI use (63.6%). Personal use is equally intense: 54.5% use AI multiple times a day, the highest of any sector.

It also records the highest severe anxiety in the study. Thirty-six point four per cent are 'very concerned' about AI affecting job security — more than double IT's 14.3% and triple the valley average of 10.0%. Sixty-three point six per cent name AI replacing journalists and content creators as a specific concern.

The response has been to move rather than resist: 68.2% are actively learning new AI skills, the second-highest rate in the study. But 13.6% say they see no need — among the highest such figures — indicating a sector splitting into two camps rather than moving together.

The trust position is the sector's most exposed feature. Net trust in AI for important decisions runs at 86.4%, the highest of any sector, in the industry where AI-generated error reaches the public most directly. Fifty-four point five per cent of the same respondents name AI-generated misinformation as a concern in their sector. Media professionals in Kathmandu recognise the risk in the abstract and extend trust anyway.

SO WHAT

Media is the sector where a professional-standards intervention would have the highest marginal value: high usage, high trust, high public exposure, and a workforce already primed to act. It is also the sector least aware of Nepal's draft AI policy (22.7%) — the lowest reading in the study.

READINESS PILLARS

PILLAR	SCORE	BASE
Literacy	**70.6**	24
Personal adoption	**79.5**	22
Workplace deployment	**81.8**	22
Workforce preparedness	**77.3**	22
Composite	**77.3**	—

WHO WAS INTERVIEWED

Journalist / reporter 43% · **Content creator** 29% · **Social media manager** 14% (base n=21)

EXPECTATIONS AND RESPONSE

Expect significant change	**50.0%**
Expect change of some kind	**81.8%**
Actively learning AI skills	**68.2%**
Net trust in AI	**86.4%**
Never use AI at work	9.1%

Sector module base n=22; some respondents did not complete the module.

EXHIBIT 21

Media & Journalism: what AI is used for, where it could add value, and what worries the sector

Multiple responses permitted; base n=22

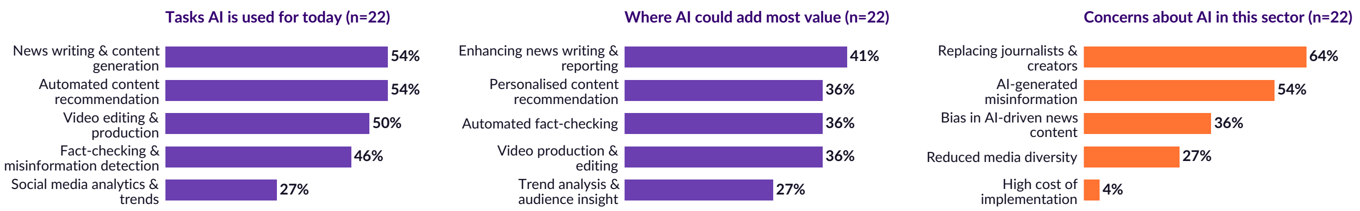

Base: Media & Journalism respondents completing the sector module, n=22. Top six options shown in each panel.

03

SECTOR DEEP DIVE · BANKING & FINANCE

Strong adoption, weak deployment against the use cases that matter most

70.1

READINESS INDEX
RANK 3 OF 7

Banking scores third of seven on the index and adopts AI readily — but almost entirely for customer-facing and advisory tasks rather than for the risk and fraud functions where the sector-specific value sits.

41	44%	95%	36%	42%	64%	39%
Respondents in sector	Report deep understanding of AI	Use AI personally	Use AI daily at work	Actively learning new AI skills	Concerned about job security	Aware of Nepal's AI policy draft

Banking & Finance ranks third with an index of 70.1, built on strong literacy (70.7) and personal adoption (75.7). Ninety-four point seven per cent use AI personally, 44.7% multiple times a day, and 36.1% use it daily at work.

The application pattern reveals the gap. The most common workplace uses are customer service chatbots (55.6%), investment advisory (44.4%) and data analysis for decisions (44.4%). **Fraud detection is used by just 5.6% of respondents**, and credit scoring and risk assessment by 27.8% — the two applications closest to the sector's own risk function.

That is not because the sector fails to see the value: asked where AI could be most useful, 27.8% name fraud detection and risk management and 33.3% name credit assessment. The gap between perceived potential and current deployment is a procurement and systems-integration gap, not an awareness gap.

Risk perception is dominated by security. Seventy-five per cent name increased cybersecurity risk as a concern — the single highest concern reading in the banking module — followed by errors in credit scoring (52.8%) and implementation cost (44.4%). Only 5.6% report no concerns.

SO WHAT

On our reading Banking is the clearest commercial opening in this study: high individual fluency, high trust, and near-zero deployment in the applications its own respondents rate most valuable. The constraint looks institutional — core-system integration and risk governance — rather than human. Budgets and purchase intent were not surveyed.

READINESS PILLARS

PILLAR	SCORE	BASE
Literacy	70.7	41
Personal adoption	75.7	38
Workplace deployment	66.0	36
Workforce preparedness	68.1	36
Composite	70.1	–

WHO WAS INTERVIEWED

Customer service rep 44% · **Manager** 31% · **Financial analyst** 11% (base n=36)

EXPECTATIONS AND RESPONSE

Expect significant change	44.4%
Expect change of some kind	86.1%
Actively learning AI skills	41.7%
Net trust in AI	63.9%
Never use AI at work	5.6%

Sector module base n=36; some respondents did not complete the module.

EXHIBIT 22

Banking & Finance: what AI is used for, where it could add value, and what worries the sector

Multiple responses permitted; base n=36

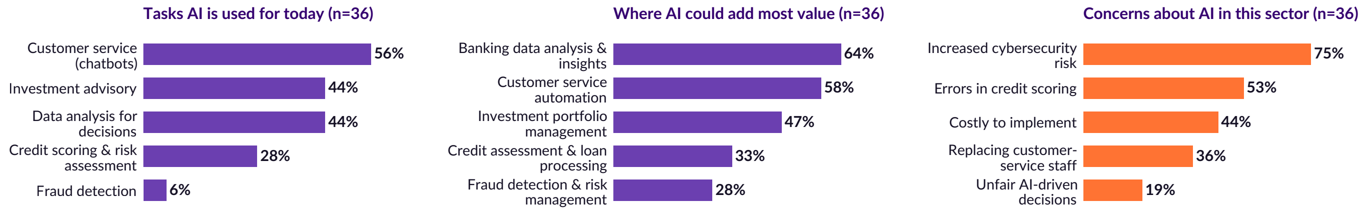

Base: Banking & Finance respondents completing the sector module, n=36. Top six options shown in each panel.

04

SECTOR DEEP DIVE · EDUCATION

High usage driven by students, thin institutional capability behind it

68.2

READINESS INDEX
RANK 4 OF 7

Education ranks fourth of seven on the index, but two-thirds of its respondents are students rather than teachers. Read as a picture of Nepal's teaching institutions rather than its learners, the sector is weaker than its score suggests — and its index score is the least comparable of the seven for that reason.

56	43%	94%	34%	34%	58%	32%
Respondents in sector	Report deep understanding of AI	Use AI personally	Use AI daily at work	Actively learning new AI skills	Concerned about job security	Aware of Nepal's AI policy draft

Education scores 68.2, above the valley average, with strong personal adoption (75.5) and workplace deployment (66.0). Thirty-four per cent use AI daily for education tasks and a further 28.0% a few times a week.

The composition of the base matters more here than in any other sector. Sixty-six per cent of Education respondents are students and 28.0% are teachers, so this page describes a mixed population of learners and educators. Splitting them changes less than expected: students and respondents with a stated non-student role report deep understanding at almost the same rate (48.5% and 47.1%), and daily use at work is near-identical (33.3% and 35.3%). **Nepali teachers in this sample are about as AI-literate as their students.** What differs is the meaning of the measure: 'workplace deployment' describes coursework for one group and teaching practice for the other, which makes this sector's index score the least comparable of the seven.

Usage is concentrated in study support: preparing study notes (72.0%) and creating lesson plans (54.0%) lead. Perceived potential is more evenly spread — adaptive learning, virtual tutoring and language translation all at 44.0%.

The concern profile is distinctive and, from an institutional standpoint, revealing. Eighty per cent worry that students will become over-dependent on AI — the joint-highest concern reading in the entire study — while only 32.0% worry about plagiarism and academic dishonesty and 30.0% about AI replacing teachers. The sector's anxiety is pedagogical rather than professional.

Preparedness is the weak pillar. Only 34.0% are actively learning new AI skills against 60.0% who 'plan to start soon' — a 26-point intention-action gap.

SO WHAT

Education is where Nepal's AI capability will either be built or lost, and the data shows it currently being neither taught nor governed — only used. Ten point seven per cent of the whole sample first learned about AI through formal education; the sector is not yet the channel it needs to be.

READINESS PILLARS

PILLAR	SCORE	BASE
Literacy	**67.1**	56
Personal adoption	**75.5**	51
Workplace deployment	**66.0**	50
Workforce preparedness	**64.0**	50
Composite	**68.2**	—

WHO WAS INTERVIEWED

Student 66% · **Teacher** 28% · **Other** 6% (base n=50)

EXPECTATIONS AND RESPONSE

Expect significant change	**48.0%**
Expect change of some kind	**82.0%**
Actively learning AI skills	**34.0%**
Net trust in AI	**56.0%**
Never use AI at work	6.0%

Sector module base n=50; some respondents did not complete the module.

EXHIBIT 23

Education: what AI is used for, where it could add value, and what worries the sector

Multiple responses permitted; base n=50

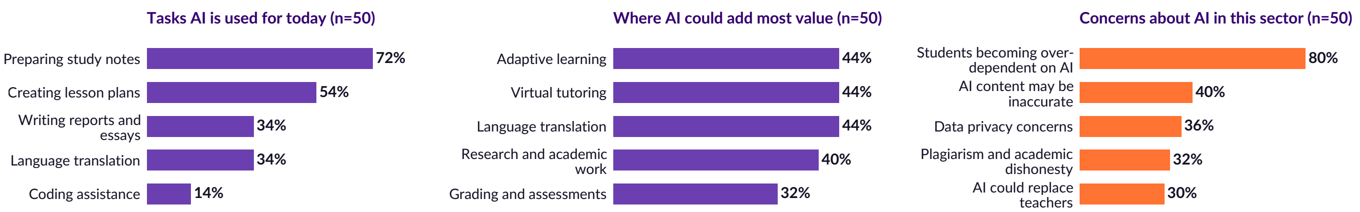

Base: Education respondents completing the sector module, n=50. Top six options shown in each panel.

05

SECTOR DEEP DIVE · TOURISM & HOSPITALITY

A literacy problem wearing the costume of complacency

Tourism records the lowest literacy score in the study and the lowest job-security concern. The second is a consequence of the first, not a sign of resilience.

57.4

READINESS INDEX
RANK 5 OF 7

36	**25%**	**75%**	**33%**	**26%**	**33%**	**48%**
Respondents in sector	Report deep understanding of AI	Use AI personally	Use AI daily at work	Actively learning new AI skills	Concerned about job security	Aware of Nepal's AI policy draft

Tourism & Hospitality scores 57.4, fifth of the seven ranked sectors. Its binding constraint is literacy (51.0) — the only sector where understanding, rather than deployment or motivation, is the weakest pillar. Only 25.0% report deep understanding of AI, against a valley figure of 38.6%.

Usage is nonetheless present and reasonably frequent: 33.3% use AI daily at work, ahead of Healthcare and Legal. But 33.3% use it only occasionally, and personal use intensity is low — 28.6% use AI multiple times a day.

Applications cluster around the customer journey: AI-powered travel booking and personalised recommendations both at 59.3%, dynamic pricing at 40.7%, chatbots at 37.0%. Where respondents see the greatest potential is different — travel security and airport automation (74.1%) and smart itinerary planning (70.4%) lead, both well ahead of current use.

Concern is the lowest in the study: 33.3% very or somewhat concerned about job security, against a valley average of 61.4%, and 29.6% explicitly not concerned at all. Given the sector's exposure to AI-driven disintermediation of travel agents and tour operators — 33.3% of the module base — this reads as under-awareness rather than security.

SO WHAT

Tourism needs foundational AI literacy before it needs tooling. It is also the sector where a low-cost, high-volume intervention would move the index furthest: raising literacy from 51.0 to the valley average of 64.5 alone would lift the sector's composite by more than three points.

READINESS PILLARS

PILLAR	SCORE	BASE
Literacy	51.0	36
Personal adoption	58.1	31
Workplace deployment	59.3	27
Workforce preparedness	61.1	27
Composite	57.4	–

WHO WAS INTERVIEWED

Travel agent / tour operator 33% · **Hotel manager / staff** 30% · **Event planner** 18% (base n=27)

EXPECTATIONS AND RESPONSE

Expect significant change	29.6%
Expect change of some kind	74.1%
Actively learning AI skills	25.9%
Net trust in AI	55.6%
Never use AI at work	3.7%

Sector module base n=27; some respondents did not complete the module.

EXHIBIT 24

Tourism & Hospitality: what AI is used for, where it could add value, and what worries the sector

Multiple responses permitted; base n=27

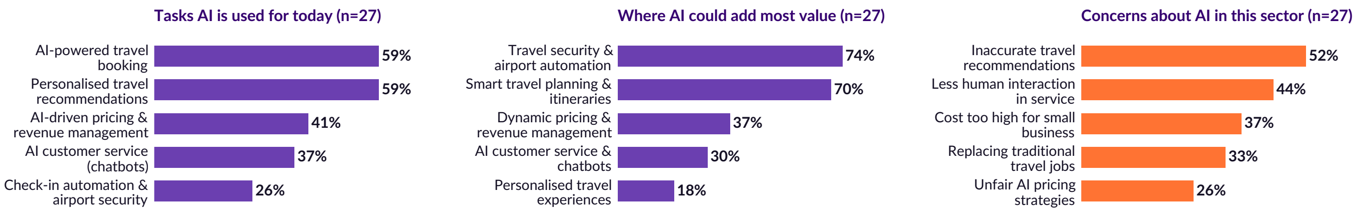

Base: Tourism & Hospitality respondents completing the sector module, n=27. Top six options shown in each panel.

06

SECTOR DEEP DIVE · HEALTHCARE

The largest sector in the study, and the least prepared for what it expects

57.1

READINESS INDEX
RANK 6 OF 7

Healthcare respondents understand AI at above-average levels, expect it to change their industry, and are the least likely of any sector to be doing anything about it. This is the widest expectation-action gap in the report.

63	41%	93%	5%	20%	63%	64%
Respondents in sector	Report deep understanding of AI	Use AI personally	Use AI daily at work	Actively learning new AI skills	Concerned about job security	Aware of Nepal's AI policy draft

Healthcare is the largest sector base in the study (n=63) and ranks sixth of seven on the index at 57.1. Its literacy score (67.6) is above the valley average of 64.5. Its workplace deployment (47.5) and workforce preparedness (53.4) are both well below.

The deployment figure is the striking one. Only 5.1% of healthcare respondents use AI daily in their work — less than half the next lowest sector with a workable base — while 32.2% use it occasionally and 25.4% a few times a month. Healthcare workers know what AI is; their working day does not include it.

Where it is used, it is used administratively. Managing patient records is the leading application at 52.5%, ahead of medical research (42.4%) and diagnosis (37.3%). Yet asked where AI could add most value, respondents name enhancing diagnostic procedures (59.3%) and improving treatment accuracy (49.2%) first. The gap between where AI is used and where it is wanted is 22 points on diagnostics.

Preparedness is the sector's critical weakness. Sixty-two point seven per cent are concerned about AI affecting job security, but only 20.3% are actively learning new skills — the lowest action rate in the study — while 66.1% 'plan to start soon'. Healthcare also records the lowest net trust in AI of any sector (37.3%), which is professionally appropriate but compounds the inertia.

The concern profile is clinically literate: patient data privacy (59.3%), AI replacing medical judgement (47.5%) and misdiagnosis (40.7%) lead. Only 6.8% report no concerns.

SO WHAT

Healthcare is the highest-priority sector for intervention in this study. It combines the largest workforce base, above-average understanding, the lowest workplace deployment, the lowest upskilling rate and the highest consequence of error. It is also the sector most aware of Nepal's draft AI policy (64.4%) — an existing channel that can be used.

READINESS PILLARS

PILLAR	SCORE	BASE
Literacy	**67.6**	63
Personal adoption	**60.0**	60
Workplace deployment	**47.5**	59
Workforce preparedness	**53.4**	59
Composite	**57.1**	–

WHO WAS INTERVIEWED

Nurse 34% · **Lab technician** 27% · **Doctor** 24% (base n=59)

EXPECTATIONS AND RESPONSE

Expect significant change	**20.3%**
Expect change of some kind	**81.4%**
Actively learning AI skills	**20.3%**
Net trust in AI	**37.3%**
Never use AI at work	8.5%

Sector module base n=59; some respondents did not complete the module.

EXHIBIT 25

Healthcare: what AI is used for, where it could add value, and what worries the sector

Multiple responses permitted; base n=59

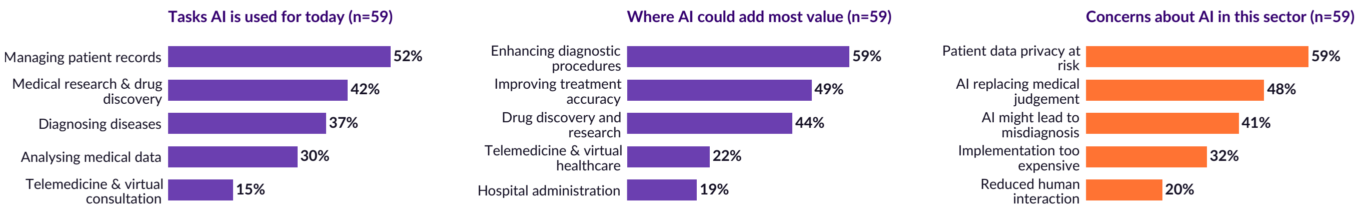

Base: Healthcare respondents completing the sector module, n=59. Top six options shown in each panel.

07

SECTOR DEEP DIVE · LEGAL

The lowest-scoring sector, held back by a single narrow use case

48.5

READINESS INDEX
RANK 7 OF 7

Legal professionals expect AI to transform their industry more than almost any other sector — and use it less than any other sector. Almost all of what use exists is one task: research.

49	**22%**	**67%**	**12%**	**28%**	**56%**	**63%**
Respondents in sector	Report deep understanding of AI	Use AI personally	Use AI daily at work	Actively learning new AI skills	Concerned about job security	Aware of Nepal's AI policy draft

Legal ranks last of the seven sectors with a workable base, with an index of 48.5 — fifteen points below the valley average and thirty-one below IT. Its personal adoption pillar (41.7) and workplace deployment pillar (38.4) are the two lowest pillar readings anywhere in the study; on adoption it sits sixteen points below the next lowest sector.

The behavioural picture is unambiguous. Sixty-two point eight per cent describe their workplace AI use as 'occasional'; only 11.6% use it daily. Two-thirds (66.7%) use AI personally against a valley figure of 87.1%, and 22.4% report deep understanding against 38.6% valley-wide.

Use is overwhelmingly concentrated in one task. Seventy-nine point one per cent use AI for conducting legal research — and nothing else comes close. Contract drafting and review registers 16.3%, legal consultation 14.0%, analysing court decisions 2.3% and predicting case outcomes 2.3%. The same concentration appears in perceived value: 72.1% name legal research and case analysis as the most useful application, with everything else below 24%.

This is not scepticism about AI's relevance. Ninety-three per cent of legal respondents expect AI to bring significant change (44.2%) or meaningful improvement (48.8%) to their industry — the highest combined expectation of any sector. The sector believes AI matters and has not built it into practice.

Concerns are appropriately professional: data privacy in legal cases (69.8%), incorrect legal advice (41.9%) and unfair AI-driven decisions (37.2%). Nine point three per cent report no concerns.

SO WHAT

Legal is the largest untapped commercial opportunity in this study — high expectation, low adoption, clear single-task beachhead already established. Any product that extends beyond research into drafting, review and document analysis is entering a market that has already accepted the category.

READINESS PILLARS

PILLAR	SCORE	BASE
Literacy	**56.9**	49
Personal adoption	**41.7**	48
Workplace deployment	**38.4**	43
Workforce preparedness	**57.0**	43
Composite	**48.5**	—

WHO WAS INTERVIEWED

Legal researcher 42% · **Lawyer** 37% · **Other** 21% (base n=43)

EXPECTATIONS AND RESPONSE

Expect significant change	**44.2%**
Expect change of some kind	**93.0%**
Actively learning AI skills	**27.9%**
Net trust in AI	**58.1%**
Never use AI at work	9.3%

Sector module base n=43; some respondents did not complete the module.

EXHIBIT 26

Legal: what AI is used for, where it could add value, and what worries the sector

Multiple responses permitted; base n=43

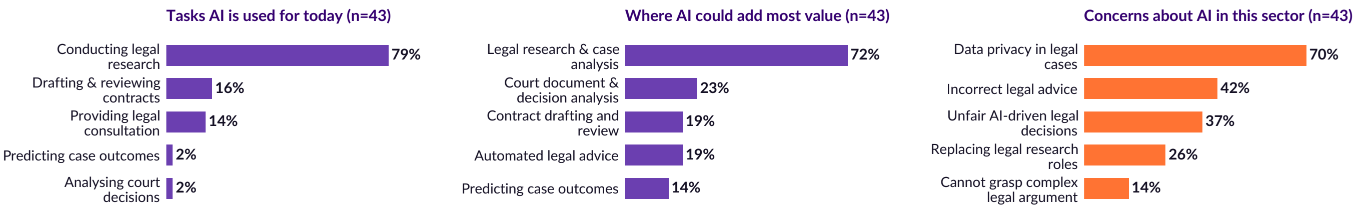

Base: Legal respondents completing the sector module, n=43. Top six options shown in each panel.

SECTOR DEEP DIVE · MANUFACTURING · EXPLORATORY ONLY, BASE OF 9

Indicative only — a base of nine

Manufacturing achieved a base of nine respondents. The figures below are reported for completeness and should not be read as representative of the sector. They are included because the directional pattern is consistent and worth testing at scale.

64.7
READINESS INDEX
NOT RANKED

9	56%	100%	0%	62%	50%	62%
Respondents in sector	Report deep understanding of AI	Use AI personally	Use AI daily at work	Actively learning new AI skills	Concerned about job security	Aware of Nepal's AI policy draft

With n=9, no percentage in this module is stable. A single respondent moves any figure by more than eleven points. The sector is retained in the index for structural completeness and is excluded from all comparative claims in this report.

The directional signals, such as they are: high self-reported literacy (68.3, the third-highest pillar reading), no daily workplace AI use at all, and a preparedness score of 75.0 driven by 62.5% actively learning.

Where AI is used, it is used for quality control and defect detection (75.0%) and production automation (37.5%). No respondent reported using AI for predictive maintenance or workplace safety monitoring — the two applications with the clearest operational return in light manufacturing.

Half name replacement of human labour as their principal concern, and a quarter report no concerns at all — the highest 'no concerns' reading in the study, and a plausible signal of low exposure rather than considered confidence.

SO WHAT

The finding for Manufacturing is methodological: Nepal's industrial sector is under-represented in AI research and should be the priority for a dedicated follow-up study with a target base of 100 or more.

READINESS PILLARS

PILLAR	SCORE	BASE
Literacy	68.3	9
Personal adoption	65.6	8
Workplace deployment	50.0	8
Workforce preparedness	75.0	8
Composite	64.7	—

WHO WAS INTERVIEWED

Technician 62% · **Factory manager** 25% · **Other** 12% (base n=8)

EXPECTATIONS AND RESPONSE

Expect significant change	0.0%
Expect change of some kind	62.5%
Actively learning AI skills	62.5%
Net trust in AI	75.0%
Never use AI at work	0.0%

Sector module base n=8; some respondents did not complete the module.

EXHIBIT 27

Manufacturing: what AI is used for, where it could add value, and what worries the sector

Multiple responses permitted; base n=8

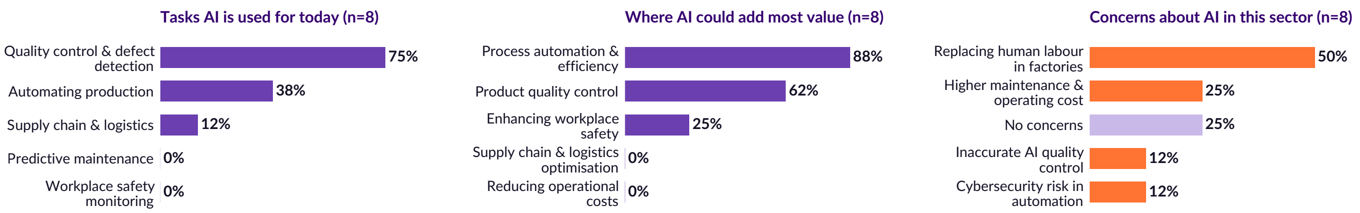

Base: Manufacturing respondents completing the sector module, n=8. Top six options shown in each panel. Base too small to support sector comparison; reported as exploratory.

08

CHAPTER EIGHT · INTERPRETATION

Four behavioural archetypes

Four segments derived from understanding and use intensity, cutting across every industry in the study. This chapter is an interpretive layer built on the survey data — purchasing behaviour, budgets and willingness to pay were not measured.

In this chapter · How the segments are built · Profile of each · Which to sell to, which to teach, which to regulate for

CHAPTER 08 · SEGMENTATION

Behaviour separates the market more sharply than industry does

Segmenting on self-assessed understanding and reported use frequency produces four groups that appear in every sector. The segments are constructed by us, not reported by respondents, and the commercial reading placed on them in this chapter is an interpretation rather than a measured market structure.

EXHIBIT 28

Four archetypes, defined by understanding and use intensity

% of n=334

Power users: deep understanding and daily use. Practical adopters: personal users at monthly-or-greater frequency who are not power users. Curious bystanders: basic understanding or occasional use. Disconnected: neither understanding nor use.

EXHIBIT 29

The archetypes diverge on preparation and policy awareness, not on demographics

% within each segment; bases n=33 to n=146

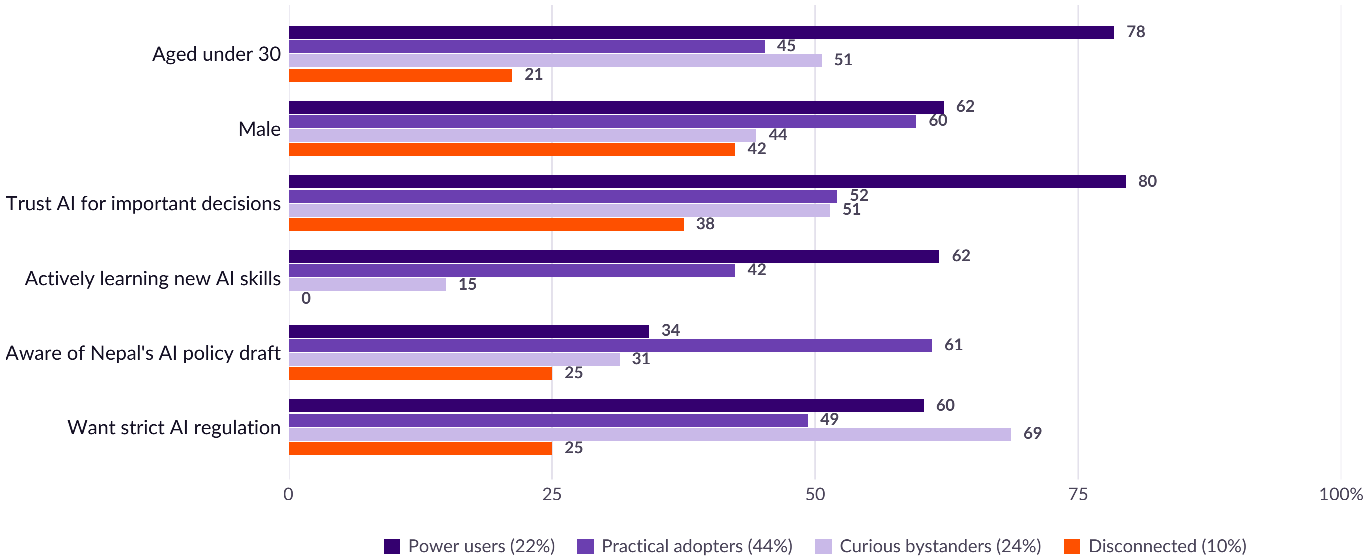

Bases: Power users n=74, Practical adopters n=146, Curious bystanders n=81, Disconnected n=33. Upskilling and policy-awareness bases are smaller where respondents did not complete the relevant module.

HOW THE SEGMENTS WERE BUILT

Two variables define the segmentation: self-assessed depth of understanding and actual frequency of use. Neither sector, age, gender nor education enters the definition — which makes the demographic differences that emerge between segments a finding rather than an assumption.

WHY THEY HOLD ACROSS SECTORS

All four segments appear in all eight sectors. Power users are most common in Education, Banking & Finance and IT; Disconnected respondents cluster in Tourism & Hospitality and Legal but appear even in IT. Behaviour, not industry, is the operative variable.

THE PATTERN THAT MATTERS MOST

Regulatory demand runs *inversely* to capability. Curious bystanders, who use AI least among the engaged, are the most likely to want strict regulation (68.6%); Practical adopters, who use it most productively, are the least (49.3%).

CHAPTER 08 · ARCHETYPE PROFILES

Who each segment is, and what each one needs

Two of the four archetypes are commercial targets. One is a policy target. One is neither, and will require a different kind of attention.

Power users · 22.2% (n=74)

Deep understanding combined with daily use. Overwhelmingly young — 78.4% are under 30, the highest of any segment — and 62.2% male. They record the highest trust in AI (79.5%) and the highest upskilling rate (61.8%). They are concentrated in Education (17), Banking & Finance (13) and IT (12).

Our reading: this group is likely to have exhausted what a general-purpose chatbot offers and is the most plausible early audience for domain tooling and advanced training. The survey did not test that. **What it does show:** only 34.2% know about Nepal's draft AI policy — the heaviest users are among the least informed about how the technology is to be regulated.

Practical adopters · 43.7% (n=146)

The largest segment. They use AI regularly but not obsessively, skew older (45.2% under 30, modal age band 30–39), and are the most highly-qualified group in the study (40.0% hold a postgraduate degree). They are also the **most policy-aware segment at 61.1%**.

Our reading: this is the group most likely to convert regular use into embedded use, given the right tools and permission. Only 42.4% are actively upskilling. We infer buying authority from age and qualification because the survey did not measure it — a gap the next wave should close.

Curious bystanders · 24.3% (n=81)

Aware and mildly engaged, but not habitual. Only 14.9% are actively learning AI skills — a quarter of the Power-user rate. They skew female relative to other segments (44.4% male, the lowest) and are concentrated in Legal (21), Healthcare (18) and Education (12).

This segment holds the strongest regulatory views in the study: **68.6% want strict AI regulation**, more than Power users (60.3%) or Practical adopters (49.3%). Distance from the technology correlates with demand for control over it.

What they need: a concrete first use case in their own professional context. Generic AI literacy training will not convert this group; a demonstration that AI does one specific thing in their job will.

Disconnected · 9.9% (n=33)

Neither understanding nor use. Older (only 21.2% under 30), less qualified (12.1% postgraduate) and **entirely absent from upskilling — 0% are actively learning**. Concentrated in Tourism & Hospitality (8), Legal (6) and IT (5).

They are also the least likely to want regulation (25.0% want strict rules) and the least policy-aware (25.0%). This is not resistance; it is disengagement.

Our reading: nothing a commercial product will reach. This is a workforce-policy responsibility, and at 10% of the sample it is a manageable one. Whether the segment grows depends on whether the pattern reflects ageing or cohort replacement — a cross-sectional survey cannot distinguish the two, and this one does not try.

THE TARGETING CONCLUSION

On our reading, the most plausible commercial audience is the 66% who are Power users or Practical adopters, and within that the 43.7% Practical adopters — qualified, regular users without depth. We stress that this is an inference from behaviour and demographics; neither budget authority nor willingness to pay was measured, and no market size is estimated anywhere in this report. For policy the priority inverts: the 34% who are Curious bystanders or Disconnected combine the highest demand for regulation with the lowest capability.

09

CHAPTER NINE

Where to act

Nine recommendations, sequenced over twenty-four months, addressed to the three parties with the ability to change the numbers in this report.

CHAPTER 09 · RECOMMENDATIONS

Nine moves, in priority order

Each recommendation is tied to a specific measured gap and, where possible, to a measurable target. None of them is an awareness campaign; awareness is at 98% and does not need help.

FOR GOVERNMENT AND POLICYMAKERS

RECOMMENDATION 01 · IMMEDIATE

Close the 45-point dissemination gap before amending the draft

91.5% want AI regulated; 46.4% know a draft policy exists. Publish a Nepali-language plain summary, distribute it through the social platforms that carried AI into the country, and target the two sectors with lowest awareness — Media & Journalism (22.7%) and Education (32.0%). **Suggested target: 75% awareness within twelve months**, set at roughly the level of the highest-performing sector today (Healthcare, 64.4%) plus a decade of normal diffusion, and measurable by repeating this single question.

MINISTRY OF COMMUNICATION AND INFORMATION TECHNOLOGY

RECOMMENDATION 02 · IMMEDIATE

Issue sector guidance for healthcare and legal practice ahead of general regulation

Healthcare respondents already use AI on patient records (52.5%) and legal respondents on case research (79.1%), with 59.3% and 69.8% respectively naming data privacy as their leading concern. These two sectors need practice notes on permissible use and data handling now — the general policy can follow. Note that only 16.9% of respondents currently support restrictions in sensitive domains, so guidance should be framed as enabling safe use rather than as prohibition.

NEPAL MEDICAL COUNCIL · NEPAL BAR ASSOCIATION · MOCIT

RECOMMENDATION 03 · 6–18 MONTHS

Build AI literacy into the tertiary curriculum, not into a standalone campaign

Only 10.7% of respondents first learned about AI through formal education, and within the education sector only 34.0% are actively building AI skills. The intervention with the longest half-life is curriculum, delivered to the cohort already in the system. **Suggested target: raise the formal-education channel share to 40% among respondents who entered the workforce after the change** — a figure that only becomes measurable in a later wave, and should be tracked by age of entry rather than by the whole sample.

UNIVERSITY GRANTS COMMISSION · MINISTRY OF EDUCATION

FOR EMPLOYERS

RECOMMENDATION 04 · IMMEDIATE

Inventory what tools and data are in use before writing a policy about them

93.6% of sector-module respondents use AI at work at least occasionally, but only 8.5% learned about AI through their employer — so most workplace AI use is self-taught and probably unrecorded. A one-page inventory of *tools, task types and data categories* converts an unknown exposure into a managed one. It should record systems and data, not individuals: an inventory that becomes employee monitoring will destroy the candour it depends on and, on this survey's evidence, drive use further underground.

ALL SECTORS · PARTICULARLY BANKING, HEALTHCARE, LEGAL

RECOMMENDATION 05 · IMMEDIATE

Set a data-handling rule for AI tools, especially where client or patient data is involved

Data privacy is the second-most-cited general concern (49.2%) and the leading sector-specific concern in Healthcare (59.3%) and Legal (69.8%). Banking's leading concern is cybersecurity risk (75.0%). Staff are already aware of the risk; what they lack is a rule telling them what may and may not be pasted into a third-party model.

BOARDS AND RISK COMMITTEES

RECOMMENDATION 06 · 6–12 MONTHS

Target training at the 30–49 cohort, where authority and fluency have come apart

Daily use falls from 48.3% among 20–29 year olds to 20.4% among 30–39 year olds and 17.1% among 40–49 year olds. The cohort that approves budgets and sets practice is the least fluent. Training aimed at "everyone" will disproportionately reach the young, who need it least. **Target: close the daily-use gap between the 20s and 30s cohorts by half.**

HR AND L&D FUNCTIONS

RECOMMENDATION 07 · 6–12 MONTHS

Convert stated intention into scheduled action

51.4% of the workforce says it plans to start learning AI skills "soon"; 39.6% actually is. Intention at this scale is an implementation problem, not a motivation problem — it responds to calendared time, assigned budget and a named owner. Healthcare (20.3% acting against 66.1% intending) and Tourism (25.9% against 70.4%) are the priorities.

SECTOR ASSOCIATIONS AND LARGE EMPLOYERS

CHAPTER 09 · RECOMMENDATIONS AND ROADMAP

For the AI industry — and the sequence in which all of this should happen

RECOMMENDATION 08 · IMMEDIATE

Look first at Legal and Banking, where expectation is high and deployment is narrow

Legal combines the highest expectation of AI-driven change in the study (93.0%) with the lowest readiness index of the ranked sectors (48.5) and a single established use case: 79.1% use AI for legal research and almost nothing else. Banking combines strong fluency and trust (63.9% net trust) with 5.6% reporting AI use for fraud detection, despite 27.8% naming it as where AI could add most value. In both sectors the category appears accepted and the application narrow. *This is a reading of reported behaviour, not a market sizing — budgets, purchase intent and willingness to pay were not surveyed.*

NEPALI AI VENDORS AND SYSTEMS INTEGRATORS

RECOMMENDATION 09 · 6–18 MONTHS

Build for the Practical adopter, not the Power user

Power users (22.2%) are already served by global general-purpose tools. Practical adopters (43.7%) are twice as numerous, older, better qualified — 40.0% hold a postgraduate degree — and use AI regularly without depth. On our reading they are the more promising audience for workflow-embedded products: design for weekly use in a specific job rather than daily use in general. The survey supports the behavioural description; the commercial conclusion is ours.

PRODUCT AND GO-TO-MARKET TEAMS

WHAT TO MEASURE IN THE NEXT WAVE

This study establishes a baseline. Four measurements would materially improve the next one:

1. **Objective capability.** Replace or supplement self-assessed understanding with a short applied-task assessment. The 38.6% deep-understanding figure is the least robust number in this report.
2. **Organisational policy.** Ask whether the respondent's employer has an AI acceptable-use policy. No question in the current instrument captures institutional response.
3. **Outcome.** Ask what AI use has changed — hours saved, tasks eliminated, errors caught. Adoption without measured outcome cannot support an economic case.
4. **Manufacturing and outside the Valley.** Boost Manufacturing to n≥100 and extend fieldwork beyond Kathmandu district, which is the entirety of the present sample.

Sequencing

The recommendations are not independent. Dissemination must precede regulation, registers must precede policies, and literacy must precede tooling in the sectors where literacy is the binding constraint. The sequence below reflects those dependencies.

HORIZON 1 — MONTHS 0–6

Close the information gaps. Low cost, no new institutions required.

- Publish plain-language Nepali summary of the draft AI policy and distribute through social channels (R1)
- Issue interim data-handling guidance for healthcare and legal practice (R2)
- Employers create AI usage registers; no policy yet, just visibility (R4)
- Boards set a minimum data rule for third-party AI tools (R5)

HORIZON 2 — MONTHS 6–18

Convert intention into capability, in the cohorts and sectors that need it.

- Targeted training programmes for the 30–49 cohort, with scheduled time and named owners (R6, R7)
- Healthcare and Tourism sector skills programmes — the two widest intention-action gaps (R7)
- AI literacy embedded into tertiary curriculum (R3)
- Vendors ship Legal and Banking workflow products beyond research and chat (R8)

HORIZON 3 — MONTHS 18–24

Institutionalise. Regulate on a base of public understanding.

- Finalise national AI policy with sectoral annexes, on a base of ≥75% public awareness (R1, R2)
- Co-regulatory framework placing duties on both state and providers, matching the 71.2% who already expect developer responsibility
- Second measurement wave, with objective capability assessment and expanded sector and geographic coverage
- Product depth for Practical adopters; sector-specific tooling as standard (R9)

APPENDIX A · METHOD

Method, data treatment and limitations

Full documentation of the analytical decisions behind every figure in this report.

Data cleaning decisions

ISSUE	TREATMENT
Aborted records	9 of 343 records contained no substantive response beyond the enumerator identifier. Removed. Analysis base n=334.
Whitespace in categorical values	Trailing spaces in the policy-awareness variable caused 137 "Yes" responses to be treated as a separate category. All string fields stripped before analysis.
Spreadsheet date corruption	Experience band "1–3 years" was stored as "01-Mar" and organisation size "10–50" as a 1950 date serial. Both restored to questionnaire categories; "201-50" restored to "201–500" and "More than 50" to "More than 500".
Municipality free text	34 distinct free-text entries normalised to local-body level. 94.0% resolve to Kathmandu Metropolitan City; six name a municipality outside Kathmandu district and four are uninterpretable and left unclassified rather than defaulted into the largest category.
Translation artefacts	Response labels in the exported file carried machine-translation artefacts from the Nepali instrument. All restored to the English questionnaire wording; the mapping is preserved in the analysis code.
Variable bases	Several questions were routed or partially completed. Every figure in this report is calculated on the base of respondents who answered that question, and the base is stated with the figure.

Statistical treatment

Because this is a purposive sample, significance testing describes associations within the sample and cannot establish that they hold in a wider population. All sixteen tests the report relies on are reported in full on the next page, with collapsed categories, Monte Carlo p-values, Cramér's V effect sizes and a false-discovery-rate correction across the whole family of tests.

Nine of the sixteen survive correction. Two claims that appeared in earlier drafts did not, and have been removed or downgraded to descriptive statements; both are identified explicitly on the following page.

The report does not weight, impute or post-stratify. Every figure is calculated on the respondents who answered that question, and that base is printed with the figure.

Index construction

Each pillar is scored as the mean of a respondent-level weight, multiplied by 100. Literacy weights: deep understanding 1.0, basic 0.5, heard-only 0.15, no idea 0. Adoption and workplace weights: daily 1.0, weekly 0.75, monthly 0.5, rarely or occasionally 0.25, never 0. Preparedness weights: actively learning 1.0, plans to start 0.5, sees no need 0.

The composite is the unweighted mean of the four pillar scores. Respondents who did not answer a pillar's underlying question are excluded from that pillar only, which is why pillar bases differ within a sector.

Limitations (1 of 2)

- **Purposive sample.** Not a probability sample; confidence intervals are not reported and results should not be projected to population totals.
- **Geographic concentration.** All respondents were recorded in Kathmandu district and 94.0% in Kathmandu Metropolitan City. The report describes urban professionals, not Nepal.
- **Youth skew.** 51.5% of the sample is under 30. Because age is the strongest predictor of AI fluency in this data, valley-level adoption figures are an upper bound.
- **Self-assessment.** Understanding, trust and preparedness are self-reported and unvalidated.

Limitations (2 of 2)

- **Small sector bases.** Manufacturing (n=9) is indicative only. Media & Journalism (n=24) supports description but not fine-grained comparison.
- **Item non-response.** 39 respondents who completed the demographic and awareness sections did not complete the perception and policy sections, reducing those bases to n=295.
- **Single point in time.** Fieldwork ran over six weeks in a fast-moving market. Tool-awareness figures in particular have a short shelf life.
- **No outcome measure.** The instrument captures use, not benefit — no figure here is a productivity estimate.

APPENDIX A · STATISTICAL TREATMENT

Effect sizes, multiple testing and index robustness

This is not a probability sample, so significance testing here describes associations *within the sample* and cannot establish that they hold in any wider population. With that limit stated, the tests below are run so as to be as reliable as the design allows.

How these tests were run

Three departures from the conventional presentation are worth stating, because each one weakens rather than strengthens the results we report.

Categories are collapsed. In the full cross-tabulations, more than half the cells in several tables had an expected count below five — age by frequency (55.6% of cells), sector by regulation preference (65.0%) — where the chi-square approximation is unreliable. Every test below is re-run on collapsed categories so that no cell is sparse.

p-values are simulated. Each p-value is computed by Monte Carlo permutation with 20,000 replications rather than from the asymptotic chi-square distribution, which makes no distributional assumption.

Multiple testing is corrected. Sixteen associations were tested. At a 0.05 threshold roughly one would be expected to appear significant by chance alone. The q column applies a Benjamini–Hochberg false-discovery-rate correction across the whole family. **Nine of the sixteen survive it.**

Cramér's V is reported alongside every test because a p-value scales with sample size and says nothing about the strength of an association. By convention V≈0.1 is weak, 0.3 moderate and 0.5 strong. **No association in this study exceeds 0.42**, and the strongest are all sector effects.

EXHIBIT 30

Every association in this study is weak to moderate

Cramér's V by association; shaded bars do not survive correction for multiple testing; n=272 to n=334

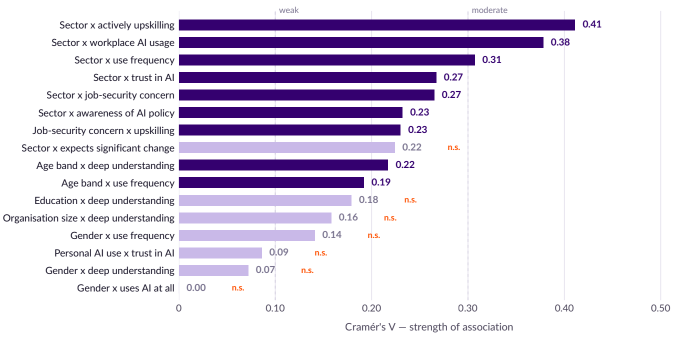

Collapsed categories, n as shown in the table opposite. "n.s." = does not survive Benjamini–Hochberg correction at 0.05.

WHAT THIS MEANS FOR THE REPORT'S CLAIMS

Two statements that appeared in earlier drafts do not survive this treatment and have been removed or qualified:

The gender difference in use frequency (q=0.079) is reported in Chapter 5 as a descriptive pattern, not an established difference.

The link between using AI and trusting it (q=0.410) does not hold, and no conclusion in this report depends on it.

Sixteen associations, tested and corrected

ASSOCIATION	N	X^2	V	P_{MC}	Q_{FDR}	SIG.
Age band x use frequency	289	21.3	0.19	<0.001	<0.001	YES
Age band x deep understanding	334	15.8	0.22	<0.001	<0.001	YES
Gender x use frequency	289	5.7	0.14	0.059	0.079	—
Gender x deep understanding	334	1.7	0.07	0.178	0.203	—
Gender x uses AI at all	310	0.0	0.00	1.000	1.000	—
Education x deep understanding	334	10.7	0.18	0.042	0.061	—
Organisation size x deep understanding	334	8.4	0.16	0.079	0.097	—
Sector x use frequency	268	50.5	0.31	<0.001	<0.001	YES
Sector x workplace AI usage	272	77.6	0.38	<0.001	<0.001	YES
Sector x trust in AI	272	38.9	0.27	<0.001	<0.001	YES
Sector x actively upskilling	272	46.0	0.41	<0.001	<0.001	YES
Sector x job-security concern	272	19.1	0.27	0.004	0.007	YES
Sector x expects significant change	272	13.6	0.22	0.032	0.051	—
Sector x awareness of AI policy	272	29.4	0.23	0.003	0.006	YES
Personal AI use x trust in AI	290	2.1	0.09	0.384	0.410	—
Job-security concern x upskilling	280	14.8	0.23	<0.001	<0.001	YES

Categories collapsed so that no expected cell count falls below five. p_{MC} = Monte Carlo permutation p-value, 20,000 replications. q_{FDR} = Benjamini–Hochberg adjusted value across all sixteen tests. Sig. = survives correction at 0.05. Sector tests exclude Manufacturing (n=9).

EXHIBIT 31

The ranking does not depend on how the index is built

Composite score under four constructions; seven ranked sectors, n=272

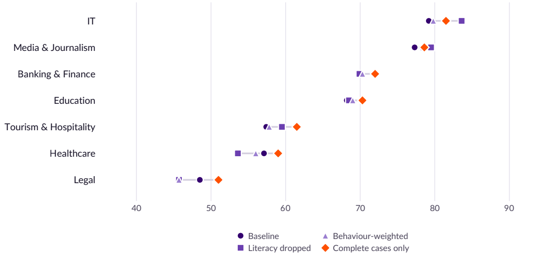

Manufacturing (n=9) excluded. Rank order is identical under all four constructions.

Is the index measuring four separate things?

Not entirely, and the report should say so. The top literacy option is worded "I understand AI well *and use it regularly*", which overlaps conceptually with personal adoption. Measured on the response data, the strongest overlap is not that pair but personal adoption and workplace deployment (Spearman ρ=0.74) — people who use AI heavily at home also use it heavily at work. Literacy correlates only moderately with adoption (ρ=0.45) and weakly with workplace deployment (ρ=0.29).

Because the pillars are not independent, the composite is best read as a summary of correlated evidence rather than as four distinct measurements. The pillar dashboard in Exhibit 9 is the more informative view; the composite is a convenience.

To test whether the ranking depends on those choices, the index was rebuilt three further ways: dropping literacy entirely; halving the weight on the two self-assessed pillars and doubling the two behavioural ones; and restricting to the 276 respondents who answered all four pillar questions. **The rank order of the seven sectors is identical under all four constructions.**

SECTOR	BASE	NO LITERACY	BEHAVIOUR-WEIGHTED	COMPLETE CASES	RANK
IT	79.2	83.6	79.8	81.5	stable
Media & Journalism	77.3	79.5	78.4	78.6	stable
Banking & Finance	70.1	69.9	70.3	72.0	stable
Education	68.2	68.5	69.0	70.3	stable
Tourism & Hospitality	57.4	59.5	57.8	61.5	stable
Healthcare	57.1	53.6	56.0	59.0	stable
Legal	48.5	45.7	45.7	51.0	stable

Complete-case variant n=276. Absolute scores move by up to five points; relative position does not move at all.

APPENDIX B · DATA TABLES

Sector summary table

All sector-level figures referenced in this report, on a single page.

SECTOR	N	INDEX	LIT.	PERS. ADOPT.	WORK. DEPLOY.	WORK. PREP.	DEEP UND.	USES AI	DAILY @WORK	LEARN-ING	JOB CONC.	NET TRUST	POLICY AWARE
IT	40	**79.2**	66.1	83.6	78.6	88.6	47.5	100.0	60.0	80.0	82.9	80.0	48.6
Media & Journalism	24	**77.3**	70.6	79.5	81.8	77.3	50.0	82.6	63.6	68.2	77.3	86.4	22.7
Banking & Finance	41	**70.1**	70.7	75.7	66.0	68.1	43.9	94.7	36.1	41.7	63.9	63.9	38.9
Education	56	**68.2**	67.1	75.5	66.0	64.0	42.9	94.1	34.0	34.0	58.0	56.0	32.0
Manufacturing (ind.)	9	**64.7**	68.3	65.6	50.0	75.0	55.6	100.0	0.0	62.5	50.0	75.0	62.5
Tourism & Hospitality	36	**57.4**	51.0	58.1	59.3	61.1	25.0	75.0	33.3	25.9	33.3	55.6	48.1
Healthcare	63	**57.1**	67.6	60.0	47.5	53.4	41.3	93.3	5.1	20.3	62.7	37.3	64.4
Legal	49	**48.5**	56.9	41.7	38.4	57.0	22.4	66.7	11.6	27.9	55.8	58.1	62.8
Valley average	334	**63.9**	64.5	66.0	59.6	65.4	38.6	87.1	29.3	39.6	61.4	58.3	46.4

All figures are percentages except n and the index/pillar scores (0–100 scale). Valley average row is calculated across all valid respondents, not as a mean of the eight sector values, and therefore includes the 16 respondents in the "Other" occupation category where the underlying question permits. Pillar bases vary — see individual sector pages.

READING THE TABLE

Pillar scores and the composite index are on a 0–100 scale. All other columns are percentages of the base that answered the relevant question, which differs by column and by sector.

COMPARABILITY

Sector columns are directly comparable with one another. The valley-average row is not a mean of the eight sector values — it is computed across all valid respondents, including the 16 in the "Other" occupation category.

CAUTION

Manufacturing (n=9) is reported for completeness only. Its values are unstable and should not be compared with those of other sectors or cited independently.

APPENDIX B · DATA TABLES

Valley-wide response distributions

Every general-question figure cited in this report, with its base.

Awareness, understanding and use

MEASURE	BASE	%
Heard of AI — yes	334	**97.6**
Deep understanding, uses regularly	334	**38.6**
Basic understanding, doesn't use much	334	48.5
Heard of it, doesn't understand	334	9.6
No idea what AI does	334	3.3
Uses AI personally — yes	310	**87.1**
Uses AI multiple times a day	289	31.1
Uses AI once a day	289	7.6
Uses AI a few times a week	289	26.6
Uses AI a few times a month	289	12.8
Uses AI rarely	289	21.5
Heard of ChatGPT	334	**90.7**
Heard of Google Gemini	334	56.3
Heard of Grammarly	334	26.0
Heard of DeepSeek AI	334	20.4
Heard of DALL·E	334	6.0
Heard of Midjourney	334	4.2
Heard of none of these	334	7.2

Channels, trust, regulation and policy

MEASURE	BASE	%
Learned of AI via social media	328	**87.2**
Learned of AI via news / articles	328	47.6
Learned of AI via friends / family	328	45.7
Learned of AI via school / university	328	10.7
Learned of AI via workplace training	328	**8.5**
Completely trust AI	295	17.3
Mostly trust, with human oversight	295	**41.0**
Neutral on trust	295	23.4
Slightly distrust, prefer humans	295	17.6
Do not trust AI at all	295	0.7
Want strict AI regulation	295	**55.9**
Want moderate AI regulation	295	35.6
Want minimal AI regulation	295	3.4
Want no AI regulation	295	0.7
Not sure about regulation	295	4.4
Heard of AI used for crime / harm	295	**82.7**
Aware of Nepal's draft AI policy	295	**46.4**
Not aware of the draft policy	295	41.0
Unsure about the draft policy	295	12.5

REFERENCES AND SOURCES

What this report relies on, and what it does not

Every statistic in the body of this report comes from the Sankhya AI survey itself. This page records the external material referred to in the text and, equally important, flags the comparisons the report deliberately does not make.

Primary source

Sankhya AI (2025). *AI Awareness and Industry-Specific Survey.* 343 structured face-to-face interviews, Kathmandu district, March–April 2025; 334 valid responses. Bilingual instrument (Nepali and English), 24 general questions plus a seven-question sector module. Anonymised dataset, analysis code and instrument available on request — see Appendix C.

Unless a figure is explicitly attributed elsewhere, it is calculated from this dataset.

Secondary source

Sankhya AI (2025). *AI Awareness and Industry-Specific Survey: questionnaire.* Nepali and English versions, 31 questions. Reproduced in full on request. The instrument is the source of every response category quoted in this report, including the wording of the question on Nepal's draft national AI policy and the list of institutions offered as possible regulators.

Why there is no external literature

This report cites no secondary literature, and that is a deliberate constraint rather than an omission. Every claim in the body is a description of what these 334 respondents reported. Where the text refers to Nepal's draft AI policy, to international organisations, or to the tasks respondents associate with their sector, it is reporting the wording of the survey instrument and the answers given to it — not asserting a fact about the world that would require an outside source.

Two consequences follow. The report cannot tell you whether Nepal's figures are high or low by any external standard. And it cannot tell you what the draft AI policy actually says — only how many respondents knew it existed.

Comparisons this report does not make

Three comparisons a reader might reasonably expect are absent, and the absence is deliberate.

- **No international benchmarking.** Adoption is not compared with India, Bangladesh, Sri Lanka or any advanced economy. Published figures use incompatible definitions of "use", different populations and different reference periods; a comparison drawn across them would mislead more than it informs.
- **No time series.** This is a first wave. Nothing here shows change over time, and no statement about trend or trajectory is supported.
- **No economic estimate.** There is no market sizing, no productivity estimate and no forecast. The instrument collected nothing that would support one.

THE EVIDENCE STANDARD USED IN THIS REPORT

Three things are kept visibly separate throughout:

The survey finds... a figure calculated directly from the response data, with its base stated.

We infer... a reading that goes beyond what was measured, flagged as such in the text.

Our judgement is... a recommendation, which carries no evidentiary weight of its own.

SUGGESTED CITATION

Giri, I. and Shrestha, R. N. (2025). *The Adoption Paradox: AI awareness, adoption and governance readiness across eight sectors in Kathmandu, Nepal.* Kathmandu: Sankhya AI. Machine-readable form on page 2.

APPENDIX C · ETHICS, CONSENT AND DATA

How respondents were treated, and what we will release

Several respondents described using AI in ways their employers were unaware of. The protections below are what made those answers possible, and they constrain what we can publish.

Consent

Every respondent was read a consent statement before any question was asked. It identified Sankhya AI as the organisation conducting the study, stated the purpose of the research, made clear that participation was voluntary, and confirmed that the respondent could decline any question or end the interview at any point.

Interviews proceeded only on explicit verbal consent. **Nine of 343 contacts did not proceed past this stage** and their records contain no substantive responses; these are the nine records excluded from the analysis base, as documented in Appendix A.

No incentive, payment or gift was offered for participation.

Enumerator conduct

The three field enumerators were trained before fieldwork on the instrument, on obtaining and recording consent, and on administering attitudinal questions without leading the respondent. Interviews were conducted in Nepali or English at the respondent's preference.

Confidentiality and personal data

The instrument collected respondent name and, optionally, a mobile number, for the sole purpose of field quality-control back-checks. Workplace name and position were collected to verify sector classification.

None of these fields are used in any analysis in this report, and all are removed from any dataset released outside Sankhya AI. No respondent, and no named employer, is identifiable anywhere in this document.

Free-text responses are reported only in aggregate. Where a verbatim answer could identify an individual or an organisation, it has been excluded rather than redacted.

A note on candour

This study asked working professionals whether they use AI at work, in a market where only 8.5% first learned about AI through their employer. Some answers may therefore describe use their organisation is unaware of. Protecting those respondents is not only an ethical obligation but a condition of the data being accurate at all.

Data availability

Sankhya AI will make the following available on request for non-commercial research, policy and journalistic use:

- The **anonymised response dataset**, with all direct identifiers removed
- The **full analysis code** — cleaning, recoding, index construction, cross-tabulation and every exhibit in this report
- The **survey instrument** in both Nepali and English
- The **verification scripts** used to check every figure in this report against the raw export

Requests should be addressed to Sankhya AI at the contact details on the inside front cover. We ask requesters to describe the intended use and to agree not to attempt re-identification of any respondent.

Corrections

If an error is identified in this report, Sankhya AI will publish a correction and reissue the affected pages. Please report suspected errors to the contact address above, citing the exhibit or page number.

REPRODUCIBILITY

Every numeric claim in this report was checked twice: once against the computed statistics file, and once by an independent script that recomputes the headline figures directly from the raw survey export without reference to the analysis pipeline. Both checks are included in the code release. Two errors identified by that process during production were corrected before publication.

About Sankhya AI

Sankhya AI is a Kathmandu-based research and applied intelligence firm. We build the evidence base that organisations in Nepal need to make decisions about artificial intelligence — through primary research, data engineering and applied deployment.

This report is based on 334 structured interviews conducted across eight sectors in Kathmandu between March and April 2025. The full anonymised dataset, analysis code and questionnaire are available on request.

CONTACT

Sankhya AI
Putalisadak, Kathmandu
Nepal

DATA AND CODE REQUESTS

The anonymised dataset, analysis code, verification scripts and bilingual questionnaire are available on request for non-commercial research, policy and journalistic use. See Appendix C.

MEDIA AND CORRECTIONS

For interviews, briefings or to report a suspected error, write to us at the address opposite, citing the exhibit or page number. Corrections are published and the affected pages reissued.